

China Internet

JPM 2026 Global China Summit Key Takeaways

- We hosted JPM 22nd Global China Summit last week in Shanghai, conducting management meetings across platform incumbents, listed AI pure-plays, and consumer internet companies spanning our coverage universe. The observation that most reframes the China AI investment debate is that AI profit pool is expanding, but economic surplus is accruing upstream, to compute owners and specialized hardware lessors, not to the model companies generating demand for that infrastructure. Inference capacity, not training scale, is the binding constraint on industry revenue growth. The market prices model-layer companies as primary AI beneficiaries while the economics we observed suggest investors should assess upstream value capture leakage before underwriting margin expansion.
- The inference bottleneck has restructured compute procurement in ways not yet reflected in consensus models, in our view. Supply has migrated from major CSPs toward specialized leasing companies, where rental costs are rising. One model company almost doubled prices yet saw only marginal margin improvement as server costs absorbed the gain. Revenue growth at the model layer is supply-constrained, not demand-constrained, a distinction with direct earnings implications.
- Model convergence is compounding this pressure. Multiple management teams acknowledged that capability gaps among leading domestic models are narrowing, shifting competition from quality toward distribution and ecosystem integration — structurally constructive for platform incumbents, more challenging for pure-plays whose valuations embed durable quality premiums. Two distinct platform strategies have crystallized: Alibaba pursuing external cloud monetization, Tencent prioritizing internal ecosystem integration, implying divergent margin profiles the market cannot evaluate with a single AI-revenue benchmark. Coding has emerged as the consensus monetization wedge, but subscription economics are already breaking down in favor of usage-based billing. No management team presented a convincing near-term path to mass consumer AI monetization, reinforcing a predominantly B2B revenue trajectory over the next 12–18 months.
- At the stock level, Tencent delivered the most significant AI narrative upgrade since 1Q26, with Hunyuan 3 adoption across 100-plus internal teams validating the ecosystem-first approach. Alibaba offered its clearest cloud margin disclosure to date, though macro commentary turned notably darker. Zhipu AI provided the sharpest illustration of upstream value capture, pricing power exists but surplus flows to lessors, ahead of the June 8 Stock Connect catalyst. MiniMax showed the strongest revenue acceleration but carries coding concentration risk and a July lock-up overhang. Baidu's candid acknowledgment of foundation model momentum loss sharpens the full-stack thesis but introduces multi-quarter execution risk.

**Head of China Equity Research
and Co-Head of Asia TMT
Research****Alex Yao** ^{AC}

(86 21) 6106 6505

alex.yao@jpmorgan.com

SAC Registration Number: S1730523020001

Daniel Chen

(86-21) 6106 6205

daniel.q.chen@jpmorgan.com

SAC Registration Number: S1730521040001

Olivia Xu

(86-21) 6106 6138

olivia.w.xu@jpmorgan.com

SAC Registration Number: S1730525060001

Nancy Liu

(86 21) 6106 6343

nancy.liu@jpmorgan.com

SAC Registration Number: S1730524090001

J.P. Morgan Securities (China) Company
Limited

Thematic Takeaways

The inference bottleneck is repricing the AI value chain in real time

The single most decision-useful observation from the Summit is that China's AI profit pool is growing but currently captured upstream, by compute owners and hardware lessors, less by the model companies generating demand for that compute.

According to management of a leading domestic foundation model developer, the company had raised pricing on its latest-generation model by roughly double, yet margins improved only marginally because server rental costs rose sharply in a single recent month with no visible ceiling. A second model company reported that API and token-plan revenue has become the majority of total revenue, up from a minority share at year-end, with text-model API gross margins meaningfully above what it characterized as industry averages, driven by throughput efficiency rather than pricing. Despite this strong unit economics profile, the company's ability to scale revenue remains constrained by inference compute availability, not demand.

We heard corroboration of this dynamic from the supply side. A global AI chip executive framed modern data centers as "token factories" where throughput converts directly to revenue and efficiency determines margins. His characterization of agentic AI systems as requiring "hundreds to thousands of times more compute" than earlier generative models reinforces the thesis that inference demand is likely to compound even as model architectures stabilize.

JPM view: the investable implication is that the market is pricing model-company equity narratives as if they are the primary beneficiaries of AI adoption, when in practice the economic surplus is flowing to compute infrastructure. This dynamic favors listed companies with vertically integrated chip-to-cloud stacks (Baidu via its proprietary AI chip subsidiary, Alibaba via T-Head/PPU) or platform incumbents that can internalize inference cost through proprietary hardware and ecosystem monetization (Tencent). Pure-play model companies (Zhipu AI, MiniMax) face the more precarious position of generating demand for compute without yet capturing proportional economics.

The compute supply chain has restructured, and the market has not caught up

According to management of multiple foundation model developers, they have shifted compute procurement away from major cloud service providers and toward specialized compute leasing companies, because the CSPs themselves are now constrained and competing for the same GPU resources they previously supplied. This is a structural change from last year, when CSPs were the primary procurement channel.

Rental cost inflation is severe and accelerating. One company described the rental price of one specific server as having risen by roughly a third in the most recent month, on top of prior months of increases. The distinction between nominal global chip production capacity and actual cards reaching China was emphasized repeatedly: supply chain disruptions involving specific server integrators were

cited as having made procurement harder and more expensive than headline production volumes imply.

Domestic chip alternatives remain a complement rather than a replacement. One company reported a domestic supplier offering only 50 cards in an entire month. The main domestic suppliers with meaningful inference volume were identified as China's leading telecom equipment maker and a domestic AI chip designer, with practical inference usability for their latest-generation chips generally estimated at Q3 2026 or later.

JPM view: model-company revenue growth over the next two to three quarters is likely constrained by compute availability more than by demand. Any easing of supply, whether through domestic chip ramp or rental price stabilization, would be a meaningful catalyst for the entire sector, benefiting Zhipu AI, MiniMax, and other model developers most directly. Conversely, continued rental inflation directly compresses the margin expansion case for model companies that lack proprietary chip supply. Cloud platforms that own proprietary chip capacity (Alibaba, Baidu) are relatively better positioned to absorb this cost pressure, though even they face constraints on available volume.

Domestic model convergence shifts competition from capability to distribution

According to management of multiple foundation model developers, the domestic competitive landscape has no generational capability gap among leading players. One described the market as not winner-take-all, with enterprise users combining models based on task-specific features (speed, long-context performance, multimodal understanding, reasoning difficulty) rather than consolidating around a single provider. A second company argued that benchmark leadership has become "easily optimized" and the meaningful test is whether paying customers choose the model in production workflows. A senior technology panel participant reinforced this from a product perspective, arguing that platform integration, user context accumulation, and workflow lock-in are becoming the durable moats.

This convergence thesis has asymmetric implications across the coverage universe. For platform incumbents, it is constructive: if the model layer is approaching "good enough" quality, then distribution, ecosystem, and data advantages become the primary competitive differentiators. One major platform demonstrated this operationally, describing how a rebuilt AI team delivered a competitive model within 88 days of full team readiness, which was then integrated across more than 100 internal product teams within weeks.

For pure-play model companies at high ARR multiples, the convergence thesis is more challenging. One such company defended premium pricing by arguing that at close to a leading US lab's capability level, discounting makes no sense because customers will simply pay slightly more for the genuine article. This logic depends on a quality gap that the broader summit evidence suggests is narrowing. We note, however, that this company's own customer feedback and retention data were described as strong. The tension between "we are the best" self-assessments and the broader "no generational gap" market structure observation is itself an important analytical signal, and the resolution likely determines the direction of AI pure-play valuations over the next 6–12 months.

JPM view: model convergence is structurally constructive for Tencent (distribution via WeChat, mini-programs, payments) and Alibaba (cloud platform, merchant ecosystem, payment rails). It is more threatening for Zhipu AI and MiniMax, whose valuations depend on durable quality premiums. Baidu sits in between: its model capability has lagged recently, but its proprietary AI chip advantage becomes more valuable in a world where the model layer commoditizes and value migrates to infrastructure.

Coding is the consensus monetization wedge, but subscription economics are failing

Across foundation model companies, cloud platforms, and internet incumbents, coding emerged as the vertical where AI revenue is actually materializing at scale. One company reported that the vast majority of token consumption is coding-related. Another described coding as its priority weakness and the area requiring the most urgent improvement. A third platform claimed its coding product leads domestic DAU rankings in the category. A Chinese AI executive positioned coding as the fastest area for top domestic models to catch up with US labs.

But the subscription model around coding is breaking. According to management across multiple sessions, coding plans are economically problematic: users exploit unlimited packages and resell unused quotas on secondary markets. At least one major platform has moved away from coding plans entirely, and management commentary across sessions suggested the industry is migrating toward usage-based token billing. One company was explicit that coding plans contribute a small minority of revenue and are more useful for reputation than profit.

JPM view: the shift from subscription to token billing has direct implications for revenue predictability and comparability. Investors evaluating model-company ARR growth should distinguish between subscription-driven and token-driven revenue, as the latter is inherently more volatile and harder to forecast, even if the unit economics are structurally better. This transition is most directly relevant for Zhipu AI and MiniMax, both of which are in the process of shifting revenue mix toward API/token billing and away from subscription-based plans.

Two diverging AI strategies are emerging among platform incumbents

The summit crystallized a strategic divergence between China's two largest Internet platforms that we expect to produce measurably different financial profiles over the next 12 months.

Alibaba is pursuing an external-first cloud strategy: sell compute and model access to everyone, including direct competitors (which were explicitly named as MaaS platform clients). According to management, the company has set a long-term ambition of US\$100 billion in external cloud revenue over approximately five years and framed the cloud business as the company's strategic core. Compute allocation prioritizes internal model training first, followed by external monetized MaaS, then infrastructure clients, and lastly the consumer AI app, a hierarchy that places external revenue high in the stack.

On the other hand, Tencent is pursuing an internal-first strategy: prioritize its own foundation model and co-designed products over external cloud customers. Management explicitly rejected making five-year AI revenue promises and framed

the strongest long-term AI advantage as ecosystem distribution (messaging, payments, mini-programs, merchants) rather than cloud revenue maximization. Compute allocation was described as structurally internal-first, with the foundation model team, co-designed products, and AI-native agent products all ranking above external cloud customers.

JPM view: these two paths imply different margin trajectories. The external-first model (Alibaba) generates more top-line AI revenue near-term but carries higher competitive-pricing and margin-compression risk. The internal-first model (Tencent) sacrifices near-term cloud growth for ecosystem lock-in and proprietary margin, but requires the model itself to reach competitive quality, a condition that management argued has now been met but that the market has not yet fully credited. Both approaches have merit; the key for investors is to evaluate each platform against the right KPI (cloud revenue growth for the former, ecosystem monetization and AI-enhanced engagement for the latter) rather than applying a single AI-revenue benchmark to both.

Consumer AI monetization remains the largest unresolved question

Across keynote speakers and corporate meetings, no management team presented a convincing near-term path to mass consumer AI monetization in China. One platform's management articulated the core problem directly: unlike traditional internet products where incremental users improve economics, in AI each additional user increases token costs materially. If token costs remain high, maximizing AI-product DAU may actually destroy value.

The proposed solutions were directional rather than proven. One platform described a future in which AI agents expand transaction volume within existing ecosystems, monetized through payment infrastructure and discovery, but acknowledged that how advertising fits into AI-native information discovery is "yet to be seen." A foundation model CEO argued that China's AI monetization should ultimately come more from the consumer market because of the country's very large user base, but that the enterprise flywheel proven in the US is unlikely to replicate in China, where the largest enterprises tend to build internally and the mid-market adopts slowly. A leading short-video platform described AI as dramatically lowering content creation costs and re-architecting recommendation systems but did not quantify how this translates to incremental monetization.

JPM view: the absence of a credible near-term consumer AI monetization path across the sector is itself an important signal. It suggests that the AI revenue ramp in China over the next 12-18 months will be predominantly B2B/API-driven, and that consumer-facing AI features are more likely to function as engagement and retention tools within existing ecosystems than as standalone revenue streams. This favors Tencent (which can embed AI into WeChat's existing transaction infrastructure) and Alibaba (which can layer AI into commerce and payment flows) over companies building consumer AI apps from scratch. For Zhipu AI and MiniMax, it reinforces the centrality of the B2B/API revenue trajectory, as consumer chatbot products are unlikely to be meaningful revenue contributors in the near term.

Company-specific takeaways

The following section examines how the thematic trends identified in Section 1, namely inference value capture, compute supply chain restructuring, model convergence, coding monetization, and the external-versus-internal AI strategy divergence, manifest at the company level. Each subsection focuses on management commentary at the JPM 2026 Global China Summit that adds granularity to these themes and sharpens company-specific investment debates.

Alibaba: Cloud margin mechanics come into focus

Key takeaway: JPM view: the margin mechanics disclosure, particularly the pricing implementation timing (with the renewal-cycle applicability accelerating the flow-through relative to a new-business-only framework), the MaaS ownership-structure economics, and the proprietary chip cost advantage, provides a more granular framework for modeling cloud profitability.

According to management, a cloud price increase took effect on April 18, 2026, applying to new contracts and renewed contracts while existing contracts retain prior pricing. This distinction matters for modeling because it implies margin improvement flows through both new business and the renewal cycle, rather than being limited to incremental volumes alone. Management described the near-term margin path as moving from approximately 9% to a low-teen rate over the next one to two quarters, driven by two near-term factors: the April pricing action and rising contribution from MaaS, which carries structurally higher margins than pure IaaS. Cost efficiency from Alibaba's proprietary PPU chip was cited as a longer-term contributor to the margin path. The longer-term margin aspiration of approximately 20% was reiterated.

On MaaS economics, management provided a breakdown by model ownership structure. On Alibaba's own Qwen models, which currently represent most of the platform usage, the company captures the full unit economics of each token served. For third-party models, the economics differ by licensing type: most third-party offerings on the platform are open-source models, for which Alibaba does not pay a developer fee and earns on the underlying compute; for closed-source third-party models, Alibaba pays a fee to the model developer, with the revenue-share varying by arrangement. Management identified several leading domestic model developers as third-party platform participants, noting that platform openness has strategic value for enterprise client attraction despite less favorable economics. Current platform usage is predominantly Qwen. Management also clarified that the externally reported MaaS ARR figure excludes internal usage from Taobao and Amap, while Ant Group (~33% owned, unconsolidated) is treated as external. This scoping detail matters because future questions about ARR quality and related-party concentration hinge on how internal traffic is categorized.

Management framed Alibaba's AI strategy as centered on proprietary model development and enterprise MaaS monetization, prioritizing B2B token revenue over consumer AI investment under binding compute constraints. Qwen is being positioned as a transaction-oriented agent, which management views as a hedge against AI-native gateway migration risk.

According to management, the company operates an explicit compute allocation hierarchy: (1) internal model training, (2) monetized MaaS and enterprise AI, (3) important infrastructure clients, and (4) the consumer Qwen app. Management indicated that the consumer app is the lowest priority for compute allocation, with growth pursued through product quality and word-of-mouth rather than heavy paid marketing, consistent with not wanting to stimulate demand for compute that cannot yet be profitably served. This hierarchy contrasts with Tencent's internal-first approach (discussed below) and reveals management's strategic view that external cloud monetization should take priority over consumer AI product development.

Tencent: Hunyuan 3 adoption exceeds expectations; consumer AI strategy takes shape

Key takeaway: Tencent's summit meetings produced the most significant upgrade in AI narrative credibility since the 1Q26 earnings call. Internal adoption of Hunyuan 3 has reached 100+ product teams, far exceeding the scope disclosed on the earnings call, and the development timeline (88 days from team assembly to established model) provides a concrete operational benchmark. Management's consumer AI monetization framework, while still exploratory, was the most specific articulation we have heard from any major Chinese platform.

According to management, Hunyuan 3's development followed a precise timeline: pre-training and post-training completed in March, internal testing began in March, and a preview version launched in April. The new team delivered the model within 88 days of having key personnel in place, a cadence metric that had not been previously disclosed. Management attributed the improvement partly to methodology changes (de-emphasizing SFT in favor of post-training) and partly to organizational rebuild. The chairman's characterization of prior-generation Hunyuan 2 as a "leaky boat" was a frank admission that the old infrastructure was inadequate and provides important context for why AI investment was deliberately held back in prior periods.

Internal adoption is materially ahead of earlier public disclosure. At the first summit session, management cited 80+ internal product teams integrated; by the second session one day later, the figure had risen to 100+, with additional teams queued. Management described co-design between model teams and product teams as a key competitive advantage: Yuanbao was the first co-design product, with gaming and QQ teams beginning co-design after Hunyuan 3 launched.

On compute allocation, management confirmed an explicit internal-first hierarchy: (1) foundation model team, (2) co-designed products, (3) AI-native and agent products, (4) external cloud customers. This was described as a structural strategic choice, not temporary triage. Management also confirmed that Q1 capex exceeded Alibaba's, a competitive positioning signal.

Management upgraded the ambition statement beyond prior public language, describing the goal as joining the "global top-10 SOTA model players", not merely leading in China. Management argued that the market underappreciates both how competitive Hunyuan 3 already is within its parameter class and how serious senior leadership is about frontier capability. However, management also

acknowledged that many leading US models operate at significantly larger scale and that larger Tencent models are still to come.

The consumer AI monetization discussion was the most substantive we have encountered from a major Chinese platform. Management explicitly rejected making five-year AI revenue promises, arguing that neither product forms nor business models are knowable on that horizon. The core economic problem was stated directly: in AI, more DAU without profitable token economics “doesn’t make sense” and could destroy value. Management described a potential crossover point where token usage rises but cost per token declines enough to make heavy consumer AI features affordable but provided no timing for this crossover.

Instead of direct subscription charging, management described a monetization path through existing ecosystem infrastructure: AI agents expanding transaction volume and monetized through payments, take-rates, merchant services, and potentially discovery and advertising. Management used a birthday-party-planning analogy to illustrate that agents need real-world vendors integrated (balloon sellers, food delivery, venue operators) and that Tencent’s mini-program ecosystem connecting tens of millions of merchants and service providers constitutes a structural advantage that pure-play AI companies cannot easily replicate.

Management also corrected a market misattribution regarding Q1 advertising margins: the margin fluctuation was driven by increased content recommendation costs, not by Hunyuan 3 rollout, which began broadly only in April.

New AI product OPEX was guided at more than Rmb36 billion for the year (more than double the prior year), covering Hunyuan, WorkBuddy, CodeBuddy, and Yuanbao, but excluding some AI spending embedded in games and advertising. Q1 spending was already running at the annualized floor of this guidance.

JPM view: the 88-day development cadence and 100+ team adoption are the most concrete evidence to date that Tencent’s AI organizational reset is producing results. The internal-first compute strategy sacrifices near-term cloud revenue growth (which management acknowledged is lagging peers for this reason) but positions the company to capture AI economics through ecosystem monetization once token costs decline sufficiently. The SOTA ambition upgrade, from competitive-in-China to global-top-10, represents a meaningful narrative shift that we expect the market to begin testing against observable model capability over the coming quarters. For the share price, the near-term catalyst remains whether Hunyuan 3 adoption translates into measurable ad tech uplift in 2Q/3Q results and whether the market credits the “from an acknowledged architectural deficit to a competitive foundation model” narrative shift with a rerate above the current level.

Baidu: Foundation model reset acknowledged; full-stack AI thesis sharpens

Key takeaway: Baidu’s summit sessions were notable for a surprisingly candid admission that the company’s foundation model effort has lost momentum over the past two years. According to management, specific chip product timelines were shared (R100 inference chip in production, R300 training + inference chip launching late 2026 to early 2027) alongside GPU cloud margin targets (mid-to-high-30s gross margin) that add granularity to the investment case.

On foundation model strategy, management made unusually candid admissions. Baidu had “the best model in China” in 2023 but slowed model training because compute and resources were diverted toward the AI search transformation. Model training was described as “inconsistent over the past two years”. The company has completed two technical team restructurings since late last year, hired new talent, and reallocated resources so the model training team can focus on training rather than supporting other applications. The target is “real tier-one/top-three” rather than permanent number-one ranking, with coding capability identified as the specific weakness requiring the most urgent investment. Management estimated meaningful improvement could come within approximately one year but it was explicit that this is not a commitment.

Management argued Baidu today resembles Google roughly 12-18 months before Gemini, TPU, and cloud momentum became visible, a conceptual analogy that frames the investment case as a full-stack AI flywheel (chips, cloud, models, applications) where the foundation model is the current weak link. If the model improves, management contends Baidu can monetize across multiple layers simultaneously: model revenue, token revenue, chip gross profit, and cloud service revenue. Management actively discouraged betting on near-term advertising rerating, describing search as a profitable cash cow whose decline should be slowed and normalized, not a growth driver.

GPU cloud margin targets were disclosed at mid-to-high-30s gross margin, with AI cloud infrastructure targeted at mid-to-high-20s. These are the first specific margin targets management has provided for these business lines.

On shareholder return, Q1 buyback execution was US\$17 million against a new three-year authorization of up to US\$5 billion, well below expectations. Management attributed the gap to geopolitical volatility disrupting execution windows and said Q2 should be higher but declined to commit to a fixed annual pace. A dividend policy is expected by year-end.

Robotaxi was presented as a hidden asset, with weekly fully driverless rides exceeding 350,000, unit economics at or beyond breakeven in Wuhan, and planned overseas expansion to Europe and the Middle East. Management referenced Waymo’s implied valuation as a comparison point, while noting that Baidu’s robotaxi business effectively receives “negative valuation” inside the consolidated entity because it remains loss-making at group level.

JPM view: the foundation model admission is important because it reframes Baidu from a company that was always behind to one that was once ahead but made a strategic allocation error, a more recoverable narrative. However, the “few quarters” timeline for model improvement means investors are unlikely to see evidence before 2H26 at earliest. For BIDU/9888, the near-term investment debate centers on whether the full-stack AI thesis (proprietary chips reducing inference cost, cloud margins expanding, model quality recovering) can gain credibility before the market loses patience. The Q1 buyback execution of US\$17 million against a US\$5 billion authorization is an incremental negative for capital-return credibility and may weigh on investor willingness to hold through the transition period. The chip product roadmap (R100 inference in production, R300 arriving late 2026 to early 2027) is important because it directly addresses the inference cost bottleneck identified in Section 1: if the proprietary chips deliver as described,

Baidu would be one of the few Chinese companies capable of structurally reducing its own inference cost, a meaningful competitive advantage as compute rental inflation persists.

Tencent Music: Ximalaya integration focus defined; ByteDance competition mechanism dissected

Key takeaway: The most decision-useful disclosures were management's clarification that the Ximalaya integration will focus on the operations most synergistic with TME's existing music and audio ecosystem, and the first explicit dissection of ByteDance's competitive mechanism (AI-generated cover uploads filling catalog gaps, supported by Douyin traffic). Management's tone was more defensive than it was in recent quarters, and the admission that revenue pressure is now broadening to advertising, not just subscriptions, represents an incremental negative.

On the Ximalaya acquisition, which closed on May 18, 2026, management indicated that the strategic focus of the integration is on Ximalaya's core long-audio and membership operations, which are the assets most synergistic with TME's existing business. These operations are profitable with gross margins above TME group average, though management described weaker operating expense and S&M discipline. The integration plan was outlined in three phases: first, integrating the content library into TME products (especially Kugou); second, improving operating efficiency; and third, pursuing longer-term cross-selling. Ximalaya is expected to be consolidated for approximately half of Q2.

On ByteDance competition, management provided the most specific description to date of the competitive mechanism. ByteDance's Qishui Music platform is using AI-generated cover versions to fill catalog gaps, effectively creating substitute content that replicates the listening experience of licensed tracks, supported by Douyin's traffic distribution advantages. This approach disproportionately affects Kugou's free and ad-supported user base. Management framed the competitive dynamic as a Tencent ecosystem versus ByteDance ecosystem contest, not merely a music-app-versus-music-app comparison.

In response, TME is escalating its legal strategy from one-by-one content takedowns to platform-level litigation, working with Tencent's legal department to pursue what management described as a landmark AI copyright precedent. This represents a structural upgrade in legal posture that, if successful, could establish industry-wide constraints on AI-generated substitute content. The outcome, however, remains uncertain, and we would not incorporate favorable resolution into estimates until legal proceedings advance.

According to management, revenue pressure is now broadening to advertising, not just subscription services – a deterioration not highlighted in prior communications. The AI remix and UGC product line was described as generating annual revenue of approximately Rmb100 million. Management also described a higher-tier membership product at approximately Rmb48 per month and an artist monetization strategy spanning concerts, merchandise, fan clubs, and ticketing.

Management's overall tone was more defensive than in recent quarters. Current guidance was described as a "candid best estimate" – language that we interpret as implying limited upside buffer and possible further downside risk. An internal

debate on Kugou's strategy, whether to widen free-entry access to retain users at the risk of cannibalizing paid subscriptions, was described as unresolved.

JPM view: TME (1698 HK/TME US) is the most defensive story in our coverage at this event. The Ximalaya acquisition is a principal structural positive, but its contribution will require execution on operating efficiency improvements to deliver material earnings accretion. The ByteDance AI-cover mechanism is structurally difficult to defend against at scale because it generates functional substitutes rather than directly infringing specific recordings, making the copyright litigation strategy the right approach but one that carries multi-quarter-to-multi-year timeline with binary outcome risk. The broadening of revenue pressure to advertising is an incremental negative not fully reflected in current estimates. We expect the stock to remain range-bound until either the litigation establishes a precedent or the Ximalaya integration produces visible earnings improvement.

Zhipu AI: Capability-led monetization, though persistently facing tight compute supply

Key takeaways: The meeting reinforced a constructive view on Zhipu as one of the few China foundation model companies with visible model quality, pricing power and enterprise demand. Feedback suggests GLM still performs better than DeepSeek V4 Pro in practical enterprise use cases, especially from internet customers on coding applications. Compute remains the main bottleneck, particularly for inference, while Zhipu's R&D compute appears better protected. The July lock-up risk also looks manageable, with only c.5% cornerstone shares involved.

GLM quality appears resilient after DS V4. Per management, the customer feedback suggests GLM remains ahead of DeepSeek V4 in coding application, with the gap more visible in back-end capability, reliability and enterprise workflow performance. DeepSeek's front-end experience is strong, but its API price discounting appears to have limited impact on GLM demand. The key takeaway is that production-grade model quality still matters more than headline benchmark or price cuts.

Compute remains tight, but Zhipu is actively alleviating the bottleneck. The main shortage is inference compute, while training capacity is more secure. Advanced chips remain the preferred option for pre-training, while domestic chips may be used more in post-training after further adaptation.

Pricing power is visible, though margin conversion remains constrained. Zhipu has raised API prices by almost 100% this year, with demand reportedly holding up well. This suggests customers are willing to pay for stronger model performance, especially in coding, agentic and reasoning workloads. However, gross margin has improved only modestly, implying much of the price increase has been absorbed by higher compute costs. Over time, margin should improve as compute pricing stabilizes, but near-term profitability will remain sensitive to GPU rental rates.

ARR growth remains strong, with API as the core revenue driver. Zhipu targets US\$1bn by year-end, and as of now, API remains the main contributor, while coding plans account for less than 10% of ARR. Overseas revenue contributes c.30% of ARR, also constrained by the compute supply. On-premise model deployment demand is also growing quickly, potentially at 50–100% this year.

Long-term differentiation comes from talent, data and focus. Zhipu remains focused on foundation models, coding, agentic workflows and reasoning. Talent stability appears strong, and the company's structure is described as Anthropic-like: clear research direction from the top with flat execution below. In the longer run, Zhipu's edge should come from algorithm talent and 2B data feedback, especially in coding where outputs can be tested and model improvement can compound with usage.

JPM view: We think the meeting reinforces Zhipu's positioning as a capability-led foundation model company with improving commercial visibility. The most important positives are resilient customer feedback after DeepSeek V4, clear ability to raise ASP, protected R&D compute access, and stronger demand in coding, agentic and reasoning workloads. Near-term margin expansion may remain constrained by inference compute inflation, but Zhipu is one of the few domestic players that can pass part of this cost pressure to customers through model pricing. We would monitor three variables from here: ARR run-rate versus the US\$1bn year-end target, gross margin conversion on fluctuating upstream computing costs, and whether enterprise usage in coding and agentic workflows continues to deepen. Overall, the meeting makes us more confident that Zhipu can capture value at the model layer, while stock volatility may still come from GPU rental inflation and fast-changing competitive headlines.

MiniMax: Accelerating ARR with a more balanced mix on LLM and multimodal demand

Key takeaways: The meeting reinforced MiniMax's positioning as an efficiency-led AI model company with rapid ARR growth, improving monetization mix, and a more aggressive roadmap toward larger models. According to management, revenue has been accelerating sequentially, with triple-digit month-over-month growth rates observed in recent months, and management expressed strong confidence in sustaining momentum through year-end. API and token plans are now the most visible growth drivers, contributing over 50% of ARR, while multimodal products and 2C subscriptions account for the balance. The next key catalyst is M3, with the nearest version expected around late May to early June targeting domestic leading in coding and agentic tasks, followed by a larger model in 3Q aimed at Opus-level capability. We think the meeting supports JPM's framework that MiniMax remains a strong cost-performance contender, with upside tied to whether M3 can convert efficiency, speed and multimodal strength into clearer model leadership.

ARR momentum remains strong, with a healthy mix. According to management, revenue has been accelerating sequentially, with triple-digit month-over-month growth rates observed in recent months, and management expressed strong confidence in sustaining momentum through year-end. A 30%-40% multimodal contribution to year-end ARR was described as reasonable, highlighting that MiniMax's revenue base is broader than text-only LLM monetization.

M3 is the next major capability catalyst, which is expected around late May to early June and targets Sonnet-level positioning, with a focus on coding, agentic capability, and differentiated multimodal understanding. Management also expects to retain MiniMax's speed advantage. The larger model is expected by late 3Q and targets Opus-level capability, with management arguing that the next major differentiation may come from being first to deliver a multi-trillion-class model.

The year-end target is to train an even larger parameter model, supported by a fast compute ramp.

Compute access is improving, with more self-controlled capacity. MiniMax started building self-controlled clusters in late 2025 through partnerships in which third parties own GPUs while MiniMax operates the clusters and uses the capacity exclusively. With international business exposure, the company is also actively enlarging compute access on a global basis.

Margin improvement is driven more by engineering than price hikes. MiniMax views token pricing as a reflection of model capability, and M2.7's blended price has already risen c.30% versus M2.5's. However, past GPM improvement was mainly driven by inference cost reduction through infra optimization, including higher KV cache hit ratio and offsite cost reduction. Management highlighted that the existing API margin is already higher than the broader industry level of c.20%, helped by 2-3x higher token throughput on similar compute resources.

Multimodal and agentic workflows are important differentiators. MiniMax sees multimodal as a large revenue opportunity beyond film and content creation, with domestic multimodal ARR growing faster than coding. Hailuo 3 is expected after the M3 release and aims to follow the LLM+DiT path, similar to Seedance 2.0, where stronger text-model understanding improves video generation. On agents, management sees the key model-layer opportunity as harness rather than standalone applications, with productivity use cases extending from coding to financial services, consulting, legal, medical research, and scientific research.

JPM view: We think the meeting reinforced MiniMax as an efficiency-led AI model company with fast ARR growth, improving compute control and differentiated multimodal monetization. The most positive points are year-to-date ARR acceleration, API/token plans rising to more than half of ARR, the M2.7 blended price up c.30%, improving API margins, and a clearer path toward self-controlled compute capacity. MiniMax's architecture advantage appears to support higher throughput, lower inference cost, and room for further pricing upside, while multimodal could become a meaningful second growth engine given higher gross margin and faster domestic ARR growth. We would monitor M3 delivery and user feedback, compute ramp execution, and whether multimodal revenue can scale beyond Hailuo subscriptions into broader use cases. Overall, the meeting makes us more constructive on MiniMax's growth visibility and margin trajectory, while the key uncertainty remains whether larger-model investment can translate into sustained customer adoption and stronger value capture.

Bilibili

AI investment focused on recommendation, video understanding, and creator efficiency

Management reiterated that the company's AI strategy remains focused on three core areas: deep video understanding, long-cycle recommendation algorithms, and creator productivity tools. Key AI initiatives include structured tagging for long-form video content, semantic integration of comments and bullet chats, and long-term interest modeling designed to reactivate older high-quality content. On the creator side, AI tools have already launched commercially and contributed to a 19% YoY increase in daily uploads. Management also disclosed that AI-driven

optimization has already improved advertising deep-conversion efficiency (CTCVR) by 25% YoY.

Advertising growth remains broad-based and sustainable

Advertising revenue increased 30% YoY in 1Q26. Management attributed the strong performance to three major drivers: improving user demographics, upgraded advertising products, and improving ad efficiency. The platform's core user base continues to turn into higher purchasing-power demographics, supporting stronger advertising demand across gaming, AI applications, consumer electronics, e-commerce, and autos. At the same time, performance advertising now accounts for more than 50% of total ad revenue, while new in-feed ad formats continue improving targeting precision and CTR performance. Importantly, management emphasized that user negative feedback rates continue declining despite monetization growth, suggesting a healthy balance between commercialization and user experience. With the current ad load still only at roughly 7%, management believes the platform retains substantial long-term monetization headroom without requiring aggressive near-term ad expansion.

Community trust and long-tail content remain core competitive advantages

Management reiterated that Bilibili does not view AI as a replacement for creators, but rather as a tool that amplifies content value and lowers production barriers. The company continues prioritizing "high-quality video" rather than emphasizing a strict long-vs-short video distinction. While short-form drama content is also growing rapidly on the platform, management highlighted that Bilibili's key advantage remains its ability to sustain engagement on older content, with many videos still generating meaningful interaction 3-6 months after release. Real user comments and bullet chats also function as a natural content-filtering mechanism, accelerating discovery of high-quality creators and strengthening community trust. In our view, the combination of creator-driven content, AI-enabled production tools, and strong monetization mechanisms creates a differentiated creator ecosystem that remains difficult for competitors to replicate.

Kanzhun

AI-driven monetization transition shifting toward outcome-based pricing

Management discussed the company's ongoing monetization upgrade strategy centered around AI-enabled recruitment efficiency and value-based pricing. The company is currently piloting new closed-loop model that with outcome-oriented pricing based on effective resume delivery. Management emphasized that the strategic opportunity in recruitment delivery services remains substantially larger than the traditional online recruitment advertising market.

AI can drive growth of recruitment service platform

Management believes AI can improve recruitment efficiency in: 1) algorithmic resumé screening to improve matching accuracy and reduce HR communication workload, and 2) large-language-model-driven communication tools to standardize initial candidate interactions and shorten response times. The company's rollout strategy remains gradual and execution-oriented. Current capabilities already include AI-assisted screening and AI-assisted communication for intent-qualified candidates, with the next phase expected to expand into AI-supported initial interviews. Over the longer term, management aims to build a fully integrated AI-

enabled recruitment delivery workflow covering candidate sourcing, screening, interviewing, and onboarding.

Recruitment demand remains resilient

From a macro demand perspective, underlying recruitment trends remain resilient. Industry-level demand is strongest in semiconductors, export-oriented manufacturing sectors (including auto sales), and AI-related industries. Meanwhile, real estate hiring declines have moderated, while service industries such as catering are gradually recovering.

Beike

Operating quality improvement driving overall margin expansion

Management highlighted that 1Q26 marked a meaningful improvement in operating quality, supported by refined organizational restructuring and more disciplined resource allocation. Despite a high base effect and softer industry conditions, the company achieved significant margin expansion through tighter cost control and operating efficiency optimization. Cash reserves remained strong in 1Q26, with management expecting operating cash flow to improve sequentially starting from 2Q26. Importantly, the restructuring strategy focused on “precision optimization” rather than broad-based employee reductions.

New home business showing signs of recovery

Management attributed the weak reported performance in the new home segment during 1Q26 primarily to a high comparison base in Jan-Feb 2025, while underlying operating trends have already begun to improve meaningfully since March. Transaction volume already turned positive YoY in Apr 2026. For 2Q26, management expect new home transaction volume to grow YoY, implying continued market share gains versus peers. The key driver of this performance is BEKE’s transition from a traditional high-commission distribution model toward a broader developer enablement and customer-centric service platform.

Rental service margin expansion sustained

Rental business continues to demonstrate substantial operating leverage improvement following product and contract model upgrades. New contracts now include annual rent adjustment rights and vacancy compensation mechanisms, while revenue recognition has shifted from gross basis to net basis accounting. As a result, the gross rental margin expanded. In our view, the key driver behind this profitability improvement is the industrialization of operational processes through AI-enabled standardization across pricing, property selection, and customer communication workflows, allowing human labor to focus increasingly on offline execution and higher-value services.

Agent role evolution becoming a long-term competitive differentiator

One of the more notable strategic discussions during the meeting centered on the evolving role of real estate agents in the AI era. Management suggested that traditional agent advantages built on information asymmetry and scripted sales techniques are rapidly diminishing as AI tools become increasingly accessible. Going forward, the core competitive differentiator for agents will shift toward higher-order capabilities such as customer need assessment, lifestyle matching, and integrated asset advisory services. In our view, this implies the industry is gradually transitioning from a labor-intensive transactional model toward a more

professionalized advisory-based framework. This evolution could structurally reshape labor economics across the brokerage industry over time.

Companies Discussed in This Report (all prices in this report as of market close on 26 May 2026, unless otherwise indicated)

Alibaba Group Holding Limited(9988.HK/HK\$127.60/OW), MiniMax Group Inc - H(0100.HK/HK\$769.00/OW), Tencent (0700)(0700.HK/HK\$439.00/OW), Tencent Music Entertainment - H(1698.HK/HK\$34.94/N), Zhipu AI(2513.HK/HK\$1345.00/OW)

Analyst Certification: The Research Analyst(s) denoted by an “AC” on the cover of this report certifies (or, where multiple Research Analysts are primarily responsible for this report, the Research Analyst denoted by an “AC” on the cover or within the document individually certifies, with respect to each security or issuer that the Research Analyst covers in this research) that: (1) all of the views expressed in this report accurately reflect the Research Analyst’s personal views about any and all of the subject securities or issuers; and (2) no part of any of the Research Analyst’s compensation was, is, or will be directly or indirectly related to the specific recommendations or views expressed by the Research Analyst(s) in this report. For all Korea-based Research Analysts listed on the front cover, if applicable, they also certify, as per KOFIA requirements, that the Research Analyst’s analysis was made in good faith and that the views reflect the Research Analyst’s own opinion, without undue influence or intervention.

All authors named within this report are Research Analysts who produce independent research unless otherwise specified. In Europe, Sector Specialists (Sales and Trading) may be shown on this report as contacts but are not authors of the report or part of the Research Department.

Important Disclosures

- **Market Maker/ Liquidity Provider:** J.P. Morgan is a market maker and/or liquidity provider in the financial instruments of/related to Tencent Music Entertainment - H, Alibaba Group Holding Limited or related entities.
- **Market Maker/ Liquidity Provider (Hong Kong):** J.P. Morgan Securities (Asia Pacific) Limited and/or J.P. Morgan Broking (Hong Kong) Limited and/or an affiliate is a market maker and/or liquidity provider in the securities of Zhipu AI, Tencent Music Entertainment - H, Alibaba Group Holding Limited, Tencent (0700), MiniMax Group Inc - H or related entities and/or warrants or options thereon, which are listed or traded on The Stock Exchange of Hong Kong Limited.
- **Manager or Co-manager:** J.P. Morgan acted as manager or co-manager in a public offering of securities or financial instruments (as such term is defined in Directive 2014/65/EU) of/for Tencent Music Entertainment - H, Alibaba Group Holding Limited, Tencent (0700) or related entities within the past 12 months.
- **Analyst Position:** An analyst on the Equity or Credit coverage team, non-fundamental analyst who may produce trade recommendations, or a member of their respective household(s) has a financial interest in the debt or equity securities of Alibaba Group Holding Limited, Tencent (0700) or related entities.
- **Beneficial Ownership (1% or more):** J.P. Morgan beneficially owns 1% or more of a class of common equity securities of Tencent Music Entertainment - H, Alibaba Group Holding Limited or related entities.
- **Client:** J.P. Morgan currently has, or had within the past 12 months, the following entity(ies) as clients: Tencent Music Entertainment - H, Alibaba Group Holding Limited, Tencent (0700), MiniMax Group Inc - H or related entities.
- **Client/Investment Banking:** J.P. Morgan currently has, or had within the past 12 months, the following entity(ies) as investment banking clients: Tencent Music Entertainment - H, Alibaba Group Holding Limited, Tencent (0700) or related entities.
- **Client/Non-Investment Banking, Securities-Related:** J.P. Morgan currently has, or had within the past 12 months, the following entity(ies) as clients, and the services provided were non-investment-banking, securities-related: Tencent Music Entertainment - H, Alibaba Group Holding Limited, Tencent (0700), MiniMax Group Inc - H or related entities.
- **Client/Non-Securities-Related:** J.P. Morgan currently has, or had within the past 12 months, the following entity(ies) as clients, and the services provided were non-securities-related: Tencent Music Entertainment - H, Alibaba Group Holding Limited, Tencent (0700) or related entities.
- **Investment Banking Compensation Received:** J.P. Morgan has received in the past 12 months compensation for investment banking services from Tencent Music Entertainment - H, Alibaba Group Holding Limited, Tencent (0700) or related entities.
- **Potential Investment Banking Compensation:** J.P. Morgan expects to receive, or intends to seek, compensation for investment banking services in the next three months from Tencent Music Entertainment - H, Alibaba Group Holding Limited, Tencent (0700), MiniMax Group Inc - H or related entities.
- **Non-Investment Banking Compensation Received:** J.P. Morgan has received compensation in the past 12 months for products or services other than investment banking from Tencent Music Entertainment - H, Alibaba Group Holding Limited, Tencent (0700), MiniMax Group Inc - H or related entities.

- **Debt Position:** J.P. Morgan may hold a position in the debt securities of Zhipu AI, Tencent Music Entertainment - H, Alibaba Group Holding Limited, Tencent (0700), MiniMax Group Inc - H or related entities, if any.

Company-Specific Disclosures: J.P. Morgan does and seeks to do business with companies covered in its research reports. As a result, investors should be aware that the firm may have a conflict of interest that could affect the objectivity of this report. Investors should consider this report as only a single factor in making their investment decision. Important disclosures, including price charts and credit opinion history tables (if applicable), are available for compendium reports and all J.P. Morgan-covered companies, and certain non-covered companies, by visiting <https://www.jpmm.com/research/disclosures>, calling 1-800-477-0406, or e-mailing research.disclosure.inquiries@jpmorgan.com with your request.

Zhipu AI (2513.HK, 2513 HK) Price Chart



Date	Rating	Price (HK\$)	Price Target (HK\$)
09-Feb-26	OW	203.20	400
27-Feb-26	OW	556.50	800
01-Apr-26	OW	693.50	950

Source: Bloomberg Finance L.P. and J.P. Morgan; price data adjusted for stock splits and dividends.
Initiated coverage Feb 09, 2026. All share prices are as of market close on the previous business day.


Tencent Music Entertainment - H (1698.HK, 1698 HK) Price Chart



Date	Rating	Price (HK\$)	Price Target (HK\$)
26-Jul-23	N	28.30	27
16-Nov-23	OW	29.30	40
22-Mar-24	OW	47.00	54
15-May-24	OW	59.20	70
15-Aug-24	OW	43.80	62
21-Mar-25	OW	56.95	64
11-May-25	OW	56.20	62
14-May-25	OW	54.50	67
14-Aug-25	OW	102.10	120
18-Mar-26	N	57.20	48
14-May-26	N	35.28	40

Source: Bloomberg Finance L.P. and J.P. Morgan; price data adjusted for stock splits and dividends.
Initiated coverage Jan 30, 2023. All share prices are as of market close on the previous business day.

Alibaba Group Holding Limited (9988.HK, 9988 HK) Price Chart



Source: Bloomberg Finance L.P. and J.P. Morgan; price data adjusted for stock splits and dividends.
Initiated coverage May 11, 2021. All share prices are as of market close on the previous business day.

Date	Rating	Price (HK\$)	Price Target (HK\$)
11-Aug-23	OW	94.30	145
17-Nov-23	OW	81.35	120
09-Jan-24	OW	70.35	105
08-Feb-24	OW	74.90	100
09-Apr-24	OW	70.40	95
16-Aug-24	OW	76.40	106
09-Oct-24	OW	104.50	120
26-Nov-24	OW	81.95	117
09-Jan-25	OW	80.90	120
21-Feb-25	OW	120.90	165
10-Jul-25	OW	102.90	135
31-Aug-25	OW	115.70	165
01-Oct-25	OW	177.00	240
09-Oct-25	OW	177.60	235
26-Nov-25	OW	157.80	225
09-Jan-26	OW	142.60	210
20-Mar-26	OW	132.00	200
08-Apr-26	OW	118.50	195
14-May-26	OW	132.80	200

Tencent (0700) (0700.HK, 700 HK) Price Chart



Source: Bloomberg Finance L.P. and J.P. Morgan; price data adjusted for stock splits and dividends.
Initiated coverage Feb 07, 2010. All share prices are as of market close on the previous business day.

Date	Rating	Price (HK\$)	Price Target (HK\$)
21-Mar-24	OW	288.80	410
15-May-24	OW	381.80	480
13-Oct-24	OW	438.80	520
20-Mar-25	OW	540.00	600
15-May-25	OW	521.00	630
16-Jul-25	OW	517.50	600
14-Aug-25	OW	586.00	685
14-Nov-25	OW	656.00	750
19-Mar-26	OW	550.50	690

MiniMax Group Inc - H (0100.HK, 100 HK) Price Chart



Date	Rating	Price (HK\$)	Price Target (HK\$)
09-Feb-26	OW	461.00	700
27-Feb-26	OW	788.50	1,000
03-Mar-26	OW	752.50	1,100

Source: Bloomberg Finance L.P. and J.P. Morgan; price data adjusted for stock splits and dividends.
Initiated coverage Feb 09, 2026. All share prices are as of market close on the previous business day.

The chart(s) show J.P. Morgan's continuing coverage of the stocks; the current analysts may or may not have covered it over the entire period.

J.P. Morgan ratings or designations: OW = Overweight, N= Neutral, UW = Underweight, NR = Not Rated

Explanation of Equity Research Ratings, Designations and Analyst(s) Coverage Universe:

J.P. Morgan uses the following rating system: Overweight (over the duration of the price target indicated in this report, we expect this stock will outperform the average total return of the stocks in the Research Analyst's, or the Research Analyst's team's, coverage universe); Neutral (over the duration of the price target indicated in this report, we expect this stock will perform in line with the average total return of the stocks in the Research Analyst's, or the Research Analyst's team's, coverage universe); and Underweight (over the duration of the price target indicated in this report, we expect this stock will underperform the average total return of the stocks in the Research Analyst's, or the Research Analyst's team's, coverage universe. NR is Not Rated. In this case, J.P. Morgan has removed the rating and, if applicable, the price target, for this stock because of either a lack of a sufficient fundamental basis or for legal, regulatory or policy reasons. The previous rating and, if applicable, the price target, no longer should be relied upon. An NR designation is not a recommendation or a rating. Some stocks under coverage have a rating but no price target; in these cases, we expect the stock will outperform/perform in line/underperform the average total return of the stocks in the Research Analyst's, or the Research Analyst's team's, coverage universe of the relevant duration of the region. In our Asia (ex-Australia and ex-India) and U.K. small- and mid-cap Equity Research, each stock's expected total return is compared to the expected total return of a benchmark country market index, not to those Research Analysts' coverage universe. If it does not appear in the Important Disclosures section of this report, the certifying Research Analyst's coverage universe can be found on J.P. Morgan's Research website, <https://www.jpmorganmarkets.com>.

Coverage Universe: Yao, Alex : Alibaba Group Holding Limited (9988.HK), Alibaba Group Holding Limited (BABA) (BABA), Alibaba Health (0241) (0241.HK), Autohome Inc (ATHM), Autohome Inc - H (2518.HK), Baidu.com (9888.HK), Baidu.com (BIDU) (BIDU), Damai Entertainment (1060.HK), DiDi Global Inc. (DIDIY), JD.Com, Inc. (9618) (9618.HK), JD.com, Inc. (JD), Kingsoft Cloud (KC), Kingsoft Cloud Holdings Ltd - H (3896.HK), Lianlian DigiTech - H (2598.HK), Maoyan Entertainment (1896.HK), Meituan (3690) (3690.HK), PDD HOLDINGS INC (PDD), Ping An Healthcare and Technology (1833) (1833.HK), Tencent (0700) (0700.HK), Tencent Music Entertainment (TME), Tencent Music Entertainment - H (1698.HK), Trip.com Group Ltd (TCOM), Trip.com Group Ltd - H (9961.HK), Vipshop (VIPS), WeRide - ADR (WRD), Weibo Corporation (WB), Weibo Corporation - H (9898.HK)

J.P. Morgan Equity Research Ratings Distribution, as of April 04, 2026

	Overweight (buy)	Neutral (hold)	Underweight (sell)
J.P. Morgan Global Equity Research Coverage*	51%	37%	12%
IB clients**	83%	79%	74%
JPMS Equity Research Coverage*	49%	39%	13%
IB clients**	94%	93%	85%

*Please note that the percentages may not add to 100% because of rounding.

**Percentage of subject companies within each of the "buy," "hold" and "sell" categories for which J.P. Morgan has provided investment banking services within the previous 12 months.

For purposes of FINRA ratings distribution rules only, our Overweight rating falls into a buy rating category; our Neutral rating falls into a hold rating category; and our Underweight rating falls into a sell rating category. Please note that stocks with an NR designation are not included in the table above. This information is current as of the end of the most recent calendar quarter.

Equity Valuation and Risks: For valuation methodology and risks associated with covered companies or price targets for covered companies, please see the most recent company-specific research report at <http://www.jpmorganmarkets.com>, contact the primary analyst or your J.P. Morgan representative, or email research.disclosure.inquiries@jpmorgan.com. For material information about the proprietary models used, please see the Summary of Financials in company-specific research reports and the Company Tearsheets, which are available to download on the company pages of our client website, <http://www.jpmorganmarkets.com>. This report also sets out within it the material underlying assumptions used.

History of Investment Recommendations:

A history of J.P. Morgan investment recommendations disseminated during the preceding 12 months can be accessed on the Research & Commentary page of <http://www.jpmorganmarkets.com> where you can also search by analyst name, sector or financial instrument.

Analysts' Compensation: The research analysts responsible for the preparation of this report receive compensation based upon various factors, including the quality and accuracy of research, client feedback, competitive factors, and overall firm revenues.

Registration of non-US Analysts: Unless otherwise noted, the non-US analysts listed on the front of this report are employees of non-US affiliates of J.P. Morgan Securities LLC, may not be registered as research analysts under FINRA rules, may not be associated persons of J.P. Morgan Securities LLC, and may not be subject to FINRA Rule 2241 or 2242 restrictions on communications with covered companies, public appearances, and trading securities held by a research analyst account.

Other Disclosures

J.P. Morgan is a marketing name for investment banking businesses of JPMorgan Chase & Co. and its subsidiaries and affiliates worldwide.

UK MIFID FICC research unbundling exemption: UK clients should refer to [UK MIFID Research Unbundling exemption](#) for details of J.P. Morgan's implementation of the FICC research exemption and guidance on relevant FICC research categorisation.

All research material made available to clients are simultaneously available on our client website, J.P. Morgan Markets, unless specifically permitted by relevant laws. Not all research content is redistributed, e-mailed or made available to third-party aggregators. For all research material available on a particular stock, please contact your sales representative.

Any long form nomenclature for references to China; Hong Kong; Taiwan; and Macau within this research material are Mainland China; Hong Kong SAR (China); Taiwan (China); and Macau SAR (China).

J.P. Morgan Research may, from time to time, write on issuers or securities targeted by economic or financial sanctions imposed or administered by the governmental authorities of the U.S., EU, UK or other relevant jurisdictions (Sanctioned Securities). Nothing in this report is intended to be read or construed as encouraging, facilitating, promoting or otherwise approving investment or dealing in such Sanctioned Securities. Clients should be aware of their own legal and compliance obligations when making investment decisions.

Any digital or crypto assets discussed in this research report are subject to a rapidly changing regulatory landscape. For relevant regulatory advisories on crypto assets, including bitcoin and ether, please see <https://www.jpmorgan.com/disclosures/cryptoasset-disclosure>.

The author(s) of this research report may not be licensed to carry on regulated activities in your jurisdiction and, if not licensed, do not hold themselves out as being able to do so.

Exchange-Traded Funds (ETFs): J.P. Morgan Securities LLC (“JPMS”) acts as authorized participant for substantially all U.S.-listed ETFs. To the extent that any ETFs are mentioned in this report, JPMS may earn commissions and transaction-based compensation in connection with the distribution of those ETF shares and may earn fees for performing other trade-related services, such as securities lending to short sellers of the ETF shares. JPMS may also perform services for the ETFs themselves, including acting as a broker or dealer to the ETFs. In addition, affiliates of JPMS may perform services for the ETFs, including trust, custodial, administration, lending, index calculation and/or maintenance and other services.

Options and Futures related research: If the information contained herein regards options- or futures-related research, such information is available only to persons who have received the proper options or futures risk disclosure documents. Please contact your J.P. Morgan Representative or visit <https://www.theocc.com/components/docs/riskstoc.pdf> for a copy of the Option Clearing Corporation's Characteristics and Risks of Standardized Options or https://www.finra.org/sites/default/files/2020-08/Security_Futures_Risk_Disclosure_Statement_2020.pdf for a copy of the Security Futures Risk Disclosure Statement.

Changes to Interbank Offered Rates (IBORs) and other benchmark rates: Certain interest rate benchmarks are, or may in the future become, subject to ongoing international, national and other regulatory guidance, reform and proposals for reform. For more information, please consult: https://www.jpmorgan.com/global/disclosures/interbank_offered_rates

Private Bank Clients: Where you are receiving research as a client of the private banking businesses offered by JPMorgan Chase & Co. and its subsidiaries (“J.P. Morgan Private Bank”), research is provided to you by J.P. Morgan Private Bank and not by any other division of J.P. Morgan, including, but not limited to, the J.P. Morgan Corporate and Investment Bank and its Global Research division.

Legal entity responsible for the production and distribution of research: The legal entity identified below the name of the Reg AC Research Analyst who authored this material is the legal entity responsible for the production of this research. Where multiple Reg AC Research Analysts authored this material with different legal entities identified below their names, these legal entities are jointly responsible for the production of this research. Where more than one legal entity is listed under an analyst's name, the first legal entity is responsible for the production unless stated otherwise. Research Analysts from various J.P. Morgan affiliates may have contributed to the production of this material but may not be licensed to carry out regulated activities in your jurisdiction (and do not hold themselves out as being able to do so). Unless otherwise stated below in the legal entity disclosures, this material has been distributed by the legal entity responsible for production, or where more than one legal entity is listed under the analyst's name, the first legal entity will be responsible for distribution. If you have any queries, please contact the relevant Research Analyst in your jurisdiction or the entity in your jurisdiction that has distributed this research material.

Legal Entities Disclosures and Country-/Region-Specific Disclosures:

Argentina: JPMorgan Chase Bank N.A Sucursal Buenos Aires is regulated by Banco Central de la República Argentina (“BCRA”- Central Bank of Argentina) and Comisión Nacional de Valores (“CNV”- Argentinian Securities Commission - ALYC y AN Integral N°51).

Australia: J.P. Morgan Securities Australia Limited (“JPMSAL”) (ABN 61 003 245 234/AFS Licence No: 238066) is regulated by the Australian Securities and Investments Commission and is a Market Participant of ASX Limited, a Clearing and Settlement Participant of ASX Clear Pty Limited and a Clearing Participant of ASX Clear (Futures) Pty Limited. This material is issued and distributed in Australia by or on behalf of JPMSAL only to “wholesale clients” (as defined in section 761G of the Corporations Act 2001). A list of all financial products covered can be found by visiting <https://www.jpmm.com/research/disclosures>. J.P. Morgan seeks to cover companies of relevance to the domestic and international investor base across all Global Industry Classification Standard (GICS) sectors, as well as across a range of market capitalisation sizes. If applicable, in the course of conducting public side due diligence on the subject company(ies), the Research Analyst team may at times perform such diligence through corporate engagements such as site visits, discussions with company representatives, management presentations, etc. Research issued by JPMSAL has been prepared in accordance with J.P. Morgan Australia's Research Independence Policy which can be found at the following link: [J.P. Morgan Australia - Research Independence Policy](#).

Brazil: Banco J.P. Morgan S.A. is regulated by the Comissão de Valores Mobiliários (CVM) and by the Central Bank of Brazil. Ombudsman J.P. Morgan: 0800-7700847 / 0800-7700810 (For Hearing Impaired) / ouvidoria.jp.morgan@jpmchase.com.

Canada: J.P. Morgan Securities Canada Inc. is a registered investment dealer, regulated by the Canadian Investment Regulatory Organization and the Ontario Securities Commission and is the participating member on Canadian exchanges. This material is distributed in Canada by or on behalf of J.P.Morgan Securities Canada Inc.

Chile: Inversiones J.P. Morgan Limitada is an unregulated entity incorporated in Chile.

China: J.P. Morgan Securities (China) Company Limited has been approved by CSRC to conduct the securities investment consultancy business.

Colombia: Banco J.P. Morgan Colombia S.A. is supervised by the Superintendencia Financiera de Colombia (SFC). Any reference in this material to products or services offered abroad by entities other than the Bank in Colombia is included exclusively for descriptive purposes. Such references do not constitute, and should not be construed as, promotional activity or the provision of financial products or services within Colombian territory, as defined under applicable Colombian regulation.

Dubai International Financial Centre (DIFC): JPMorgan Chase Bank, N.A., Dubai Branch is regulated by the Dubai Financial Services Authority (DFSA) and its registered address is Dubai International Financial Centre - The Gate, West Wing, Level 3 and 9 PO Box 506551, Dubai, UAE. This material has been distributed by JP Morgan Chase Bank, N.A., Dubai Branch to persons regarded as professional clients or market counterparties as defined under the DFSA rules.

European Economic Area (EEA): Unless specified to the contrary, research is distributed in the EEA by J.P. Morgan SE ("JPM SE"), which is authorised as a credit institution by the Federal Financial Supervisory Authority (Bundesanstalt für Finanzdienstleistungsaufsicht, BaFin) and jointly supervised by the BaFin, the German Central Bank (Deutsche Bundesbank) and the European Central Bank (ECB). JPM SE is a company headquartered in Frankfurt with registered address at TaunusTurm, Taunustor 1, Frankfurt am Main, 60310, Germany. The material has been distributed in the EEA to persons regarded as professional investors (or equivalent) pursuant to Art. 4 para. 1 no. 10 and Annex II of MiFID II and its respective implementation in their home jurisdictions ("EEA professional investors"). This material must not be acted on or relied on by persons who are not EEA professional investors. Any investment or investment activity to which this material relates is only available to EEA relevant persons and will be engaged in only with EEA relevant persons.

Hong Kong: J.P. Morgan Securities (Asia Pacific) Limited (CE number AAJ321) is regulated by the Hong Kong Monetary Authority and the Securities and Futures Commission in Hong Kong, and J.P. Morgan Broking (Hong Kong) Limited (CE number AAB027) is regulated by the Securities and Futures Commission in Hong Kong. JP Morgan Chase Bank, N.A., Hong Kong Branch (CE Number AAL996) is regulated by the Hong Kong Monetary Authority and the Securities and Futures Commission, is organized under the laws of the United States with limited liability. Where the distribution of this material is a regulated activity in Hong Kong, the material is distributed in Hong Kong by or through J.P. Morgan Securities (Asia Pacific) Limited and/or J.P. Morgan Broking (Hong Kong) Limited.

India: J.P. Morgan India Private Limited (Corporate Identity Number - U67120MH1992FTC068724), having its registered office at J.P. Morgan Tower, Off. C.S.T. Road, Kalina, Santacruz - East, Mumbai – 400098, is registered with the Securities and Exchange Board of India (SEBI) as a 'Research Analyst' having registration number INH000001873. J.P. Morgan India Private Limited is also registered with SEBI as a member of the National Stock Exchange of India Limited and the Bombay Stock Exchange Limited (SEBI Registration Number – INZ000239730) and as a Merchant Banker (SEBI Registration Number - MB/INM000002970). Telephone: 91-22-6157 3000, Facsimile: 91-22-6157 3990 and Website: <http://www.jpimipl.com>. JPMorgan Chase Bank, N.A. - Mumbai Branch is licensed by the Reserve Bank of India (RBI) (Licence No. 53/ Licence No. BY.4/94; SEBI - IN/CUS/014/ CDSL : IN-DP-CDSL-444-2008/ IN-DP-NSDL-285-2008/ INBI00000984/ INE231311239) as a Scheduled Commercial Bank in India, which is its primary license allowing it to carry on Banking business in India and other activities, which a Bank branch in India are permitted to undertake. For non-local research material, this material is not distributed in India by J.P. Morgan India Private Limited. Compliance Officer: Ashutosh Sharma; ashutosh.j.sharma@jpmchase.com; +912261575002. Grievance Officer: Ramprasadh K, jpimipl.research.feedback@jpmorgan.com; +912261573000. Registration granted by SEBI and certification from NISM in no way guarantee performance of the intermediary or provide any assurance of returns to investors. Please visit [Terms and Conditions and Most Important Terms and Conditions \(MITC\)](#). The annual Compliance audit report is available at <http://www.jpimipl.com/#research>.

Indonesia: PT J.P. Morgan Sekuritas Indonesia is a member of the Indonesia Stock Exchange and is registered and supervised by the Otoritas Jasa Keuangan (OJK).

Korea: J.P. Morgan Securities (Far East) Limited, Seoul Branch, is a member of the Korea Exchange (KRX). JPMorgan Chase Bank, N.A., Seoul Branch, is licensed as a branch office of foreign bank (JPMorgan Chase Bank, N.A.) in Korea. Both entities are regulated by the Financial Services Commission (FSC) and the Financial Supervisory Service (FSS). For non-macro research material, the material is distributed in Korea by or through J.P. Morgan Securities (Far East) Limited, Seoul Branch.

Japan: JPMorgan Securities Japan Co., Ltd. and JPMorgan Chase Bank, N.A., Tokyo Branch are regulated by the Financial Services Agency in Japan.

Malaysia: This material is issued and distributed in Malaysia by JPMorgan Securities (Malaysia) Sdn Bhd (18146-X), which is a Participating Organization of Bursa Malaysia Berhad and holds a Capital Markets Services License issued by the Securities Commission in Malaysia.

Mexico: J.P. Morgan Casa de Bolsa, S.A. de C.V., J.P. Morgan Grupo Financiero is member of the Mexican Stock Exchange ("Bolsa Mexicana de Valores") and the Institutional Stock Exchange ("Bolsa Institucional de Valores"), and it is authorized to act as a broker dealer by the National Banking and Securities Exchange Commission ("Comisión Nacional Bancaria y de Valores").

New Zealand: This material is issued and distributed by JPMSAL in New Zealand only to "wholesale clients" (as defined in the Financial Markets Conduct Act 2013). JPMSAL is registered as a Financial Service Provider under the Financial Service providers (Registration and Dispute Resolution) Act of 2008.

Philippines: J.P. Morgan Securities Philippines Inc. is a Trading Participant of the Philippine Stock Exchange and a member of the Securities Clearing Corporation of the Philippines and the Securities Investor Protection Fund. It is regulated by the Securities and Exchange Commission.

Singapore: This material is issued and distributed in Singapore by or through J.P. Morgan Securities Singapore Private Limited (JPMS) [MDDI (P) 057/08/2025 and Co. Reg. No.: 199405335R], which is a member of the Singapore Exchange Securities Trading Limited, and/or JPMorgan Chase Bank, N.A., Singapore branch (JPMCB Singapore), both of which are regulated by the Monetary Authority of Singapore. This material is issued and distributed in Singapore only to accredited investors, expert investors and institutional investors, as defined in Section 4A of the Securities and Futures Act, Cap. 289 (SFA). This material is not intended to be issued or distributed to any retail investors or any other investors that do not fall into the classes of "accredited investors," "expert investors" or "institutional

investors,” as defined under Section 4A of the SFA. Recipients of this material in Singapore are to contact JPMSS or JPMCB Singapore in respect of any matters arising from, or in connection with, the material.

South Africa: J.P. Morgan Equities South Africa Proprietary Limited and JPMorgan Chase Bank, N.A., Johannesburg Branch are members of the Johannesburg Securities Exchange and are regulated by the Financial Services Conduct Authority (FSCA).

Taiwan: J.P. Morgan Securities (Taiwan) Limited is a participant of the Taiwan Stock Exchange (company-type) and regulated by the Taiwan Securities and Futures Bureau. Material relating to equity securities is issued and distributed in Taiwan by J.P. Morgan Securities (Taiwan) Limited, subject to the license scope and the applicable laws and the regulations in Taiwan. **To the extent that J.P. Morgan Securities (Taiwan) Limited produces research materials on securities not listed on the Taiwan Stock Exchange or Taipei Exchange (“Non-Taiwan Listed Securities”), these materials shall not constitute securities recommendations for the purpose of applicable Taiwan regulations, and, for the avoidance of doubt, J.P. Morgan Securities (Taiwan) Limited does not act as broker for Non-Taiwan Listed Securities.** According to Paragraph 2, Article 7-1 of Operational Regulations Governing Securities Firms Recommending Trades in Securities to Customers (as amended or supplemented) and/or other applicable laws or regulations, please note that the recipient of this material is not permitted to engage in any activities in connection with the material that may give rise to conflicts of interests, unless otherwise disclosed in the “Important Disclosures” in this material.

Thailand: This material is issued and distributed in Thailand by JPMorgan Securities (Thailand) Ltd., which is a member of the Stock Exchange of Thailand and is regulated by the Ministry of Finance and the Securities and Exchange Commission. The registered address is 548 One City Center Building, 50th Floor, Ploenchit Road, Lumpini, Pathum Wan, Bangkok 10330.

UK: Research is produced in the UK by J.P. Morgan Securities plc (“JPMS plc”) which is a member of the London Stock Exchange and is authorised by the Prudential Regulation Authority and regulated by the Financial Conduct Authority and the Prudential Regulation Authority or J.P. Morgan Markets Limited (“JPMML Ltd”) which is authorised and regulated by the Financial Conduct Authority. Unless specified to the contrary, this material is distributed in the UK by JPMS plc and is directed in the UK only to: (a) persons having professional experience in matters relating to investments falling within article 19(5) of the Financial Services and Markets Act 2000 (Financial Promotion) (Order) 2005 (“the FPO”); (b) persons outlined in article 49 of the FPO (high net worth companies, unincorporated associations or partnerships, the trustees of high value trusts, etc.); or (c) any persons to whom this communication may otherwise lawfully be made; all such persons being referred to as “UK relevant persons”. This material must not be acted on or relied on by persons who are not UK relevant persons. Any investment or investment activity to which this material relates is only available to UK relevant persons and will be engaged in only with UK relevant persons. A description of J.P. Morgan EMEA’s policy for prevention and avoidance of conflicts of interest related to the production of Research can be found at the following link: [J.P. Morgan EMEA - Research Independence Policy](#).

U.S.: J.P. Morgan Securities LLC (“JPMS”) is a member of the NYSE, FINRA, SIPC, and the NFA. JPMorgan Chase Bank, N.A. is a member of the FDIC. Material published by non-U.S. affiliates is distributed in the U.S. by JPMS who accepts responsibility for its content.

General: Additional information is available upon request. The information in this material has been obtained from sources believed to be reliable. While all reasonable care has been taken to ensure that the facts stated in this material are accurate and that the forecasts, opinions and expectations contained herein are fair and reasonable, JPMorgan Chase & Co. or its affiliates and/or subsidiaries (collectively J.P. Morgan) make no representations or warranties whatsoever to the completeness or accuracy of the material provided, except with respect to any disclosures relative to J.P. Morgan and the Research Analyst’s involvement with the issuer that is the subject of the material. Accordingly, no reliance should be placed on the accuracy, fairness or completeness of the information contained in this material. There may be certain discrepancies with data and/or limited content in this material as a result of calculations, adjustments, translations to different languages, and/or local regulatory restrictions, as applicable. These discrepancies should not impact the overall investment analysis, views and/or recommendations of the subject company(ies) that may be discussed in the material. Artificial intelligence tools may have been used in the preparation of this material, including assisting in data analysis, pattern recognition, and content drafting for research material. J.P. Morgan accepts no liability whatsoever for any loss arising from any use of this material or its contents, and neither J.P. Morgan nor any of its respective directors, officers or employees, shall be in any way responsible for the contents hereof, apart from the liabilities and responsibilities that may be imposed on them by the relevant regulatory authority in the jurisdiction in question, or the regulatory regime thereunder. Opinions, forecasts or projections contained in this material represent J.P. Morgan’s current opinions or judgment as of the date of the material only and are therefore subject to change without notice. Periodic updates may be provided on companies/industries based on company-specific developments or announcements, market conditions or any other publicly available information. There can be no assurance that future results or events will be consistent with any such opinions, forecasts or projections, which represent only one possible outcome. Furthermore, such opinions, forecasts or projections are subject to certain risks, uncertainties and assumptions that have not been verified, and future actual results or events could differ materially. The value of, or income from, any investments referred to in this material may fluctuate and/or be affected by changes in exchange rates. All pricing is indicative as of the close of market for the securities discussed, unless otherwise stated. Past performance is not indicative of future results. Accordingly, investors may receive back less than originally invested. This material is not intended as an offer or solicitation for the purchase or sale of any financial instrument. The opinions and recommendations herein do not take into account individual client circumstances, objectives, or needs and are not intended as recommendations of particular securities, financial instruments or strategies to particular clients. This material may include views on structured securities, options, futures and other derivatives. These are complex instruments, may involve a high degree of risk and may be appropriate investments only for sophisticated investors who are capable of understanding and assuming the risks involved. The recipients of this material must make their own

independent decisions regarding any securities or financial instruments mentioned herein and should seek advice from such independent financial, legal, tax or other adviser as they deem necessary. J.P. Morgan may trade as a principal on the basis of the Research Analysts' views and research, and it may also engage in transactions for its own account or for its clients' accounts in a manner inconsistent with the views taken in this material, and J.P. Morgan is under no obligation to ensure that such other communication is brought to the attention of any recipient of this material. Others within J.P. Morgan, including Strategists, Sales staff and other Research Analysts, may take views that are inconsistent with those taken in this material. Employees of J.P. Morgan not involved in the preparation of this material may have investments in the securities (or derivatives of such securities) mentioned in this material and may trade them in ways different from those discussed in this material. This material is not an advertisement for or marketing of any issuer, its products or services, or its securities in any jurisdiction.

Confidentiality and Security Notice: This transmission may contain information that is privileged, confidential, legally privileged, and/or exempt from disclosure under applicable law. If you are not the intended recipient, you are hereby notified that any disclosure, copying, distribution, or use of the information contained herein (including any reliance thereon) is STRICTLY PROHIBITED. Although this transmission and any attachments are believed to be free of any virus or other defect that might affect any computer system into which it is received and opened, it is the responsibility of the recipient to ensure that it is virus free and no responsibility is accepted by JPMorgan Chase & Co., its subsidiaries and affiliates, as applicable, for any loss or damage arising in any way from its use. If you received this transmission in error, please immediately contact the sender and destroy the material in its entirety, whether in electronic or hard copy format. This message is subject to electronic monitoring: <https://www.jpmorgan.com/disclosures/email>

MSCI: Certain information herein ("Information") is reproduced by permission of MSCI Inc., its affiliates and information providers ("MSCI") ©2026. No reproduction or dissemination of the Information is permitted without an appropriate license. MSCI MAKES NO EXPRESS OR IMPLIED WARRANTIES (INCLUDING MERCHANTABILITY OR FITNESS) AS TO THE INFORMATION AND DISCLAIMS ALL LIABILITY TO THE EXTENT PERMITTED BY LAW. No Information constitutes investment advice, except for any applicable Information from MSCI ESG Research. Subject also to [msci.com/disclaimer](https://www.msci.com/disclaimer)

Sustainalytics: Certain information, data, analyses and opinions contained herein are reproduced by permission of Sustainalytics and: (1) includes the proprietary information of Sustainalytics; (2) may not be copied or redistributed except as specifically authorized; (3) do not constitute investment advice nor an endorsement of any product or project; (4) are provided solely for informational purposes; and (5) are not warranted to be complete, accurate or timely. Sustainalytics is not responsible for any trading decisions, damages or other losses related to it or its use. The use of the data is subject to conditions available at <https://www.sustainalytics.com/legal-disclaimers>. ©2026 Sustainalytics. All Rights Reserved.

"Other Disclosures" last revised May 16, 2026.

Copyright 2026 J.P. Morgan Securities (China) Company Limited. All rights reserved. This material or any portion hereof may not be reprinted, sold or redistributed without the written consent of J.P. Morgan. It is strictly prohibited to use or share without prior written consent from J.P. Morgan any research material received from J.P. Morgan or an authorized third-party ("J.P. Morgan Data") in any third-party artificial intelligence ("AI") systems or models when such J.P. Morgan Data is accessible by a third-party.