

AI Series #63: China CSP Focus Areas in Earnings; Recent Models Post Surprises

Highlights (1) **fast model iterations and cost efficiencies** are the trends as seen from China and US models ([link](#)); (2) Continued breakthrough in **model intelligence, agent and coding capabilities** from **K3, DS-V4 Flash, Qwen3.8-Max**; (3) **Token Expenditure Index**. Week of 26 Jul continued to be soft. It reflects continued appeal of cost-efficient models ([Exhibit 4](#)); (4) **China CSP expectations** in upcoming results post US earnings.


In this report, we provide investors with over 50 data points on model pricing by companies, pricing comparison between China and US, model intelligence from Artificial Analysis, token consumption in OpenRouter, Token Expenditure Index trends and other industry data for analysis.

What do we expect from China CSP in coming results. Highlights: **BABA**: (1) We estimate **AliCloud revenue** accelerates to 45% YoY and **further accelerates in Sep-Q**; (2) We expect **AI revenue contribution** as % of external revenue to increase from about 30% in Jun-Q to 50% in a year time; (3) AliCloud **margin improves sequentially** in Jun-Q, thanks to rising AI revenue contribution; (4) ARR is expected to reach USD1.5bn in Jun-Q. We expect new updates on full year outlook at USD4.5bn by year end; (5) T-Head outlook which is important to full stack capabilities; (6) We expect **ecommerce EBITA (domestic + overseas)** to be flattish in Jun-Q despite macro-uncertainties and supports investments in AI and cloud; **Baidu**: (1) we expect Baidu AI Cloud Infra revenue to increase by 55% YoY (46% YoY including AI app); (2) AI investments to step up in 2H (e.g. models). We estimate Baidu Core non-GAAP operating profit to decline YoY in 2026; **Tencent**: (1) we expect WorkBuddy to be the focus in upcoming earnings results. According to Analysis (Jun-26), no of desktop monthly visits reached about 21m, ahead of TRADE IDE (about 12.8m) and QoderWork (about 7.9m); (2) Updates about the testing of Weixin AI; (3) Hy4 progress; (4) On Capex, we estimate it to notably increase in 2H vs 1H; **Kingsoft Cloud** (1) we estimate revenue to increase by 28% YoY considering supply constraint issues in Apr and May before returning to normal in Jun; (2) we expect revenue acceleration in 3Q; (3) margin improvement in 2Q vs 1Q. We expect market to focus on discussions (based on US earnings results) which include (1) Capex; (2) FCF and ROIC; (3) margin outlook; (4) importance of own models and model competition.

Fast model iterations and cost efficiencies. Post the release of Kimi K3, DS V4 Flash, Qwen3.8-Max and H3, (1) intensifying competition is expected with more 3T models to come in China ([link](#)); (2) an increasing emphasis on cost efficiencies as can be seen from the trends in US ([link](#)) and China ([Exhibit 2](#)). For DS V4-Flash, the cost per task is USD0.03 only, which is materially lower than overseas and domestic models ([Exhibit 15](#)). For H3, it is open weight and priced about half the price of peers at 768p and 1/3 at 2K resolution; (3) full stack capabilities is important. The readthrough from US earnings results shows (1) own models can have additional control over costs (e.g. applications) and cost-effective solutions for customers; (2) one model cannot dominate the market. Developers use different models to build applications. There will be at least half a dozen models which are comparably as good as each other.

Continued overleaf...

Exhibit 1 - Artificial Intelligence vs Blended Price



Source: Artificial Analysis, Jefferies

Exhibit 2 - Flagship models by companies

Model	Total parameter No. of Active parameters	AA Intelligence Index Score	Chinese Input Price (USD)	Chinese Output Price (USD)	Qwen 3.8 Input Price (USD)	Qwen 3.8 Output Price (USD)
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.80	0.00	11.80
Qwen 3.8-Max	2.2B	508	12.7	3.8		

Silicon Data LLM Token Expenditure Index: The week of 26 Jul continues to show softness.

The index ([Exhibit 4](#)) is weighted by where usage is concentrated (e.g. expenditure goes up when demand shifts towards premium models). On the other hand, media reports (e.g. *The Information*, *Economic Observer*) highlight tech and internet companies set upper limit on token consumption among employees. According to our checks, the index ranged from 1.36 to 1.45 the week of 26 Jul (vs 1.47 to 1.58 the week of 19 Jul), vs 2.04 on 31 May. Based on Artificial Analysis, China models DeepSeek, Qwen, Kimi, MiniMax and Zhipu are highly cost-efficient vs US models and China API cost is a fraction vs US models ([Exhibit 1](#)), thanks to former's refined architecture such as MoE (Mixture of Experts), attention mechanism (Group-Query Attention GQA, sparse attention) and MFU (Model FLOPs Utilization).

OpenRouter: Token consumption came down 2% post promotions in prior weeks. Based on analysis of the week of 27 Jul (vs week of 20 Jul) ([Exhibit 3](#)), number of weekly tokens consumed declined by 2.1% to 56.8tn. In terms of ranking, DeepSeek V4 Flash ranked no 1, followed by Xiaomi MiMo-V2.5 and Tencent Hy3.

BAT, KC, AI labs and Kling are Beneficiaries on Rising AI Demand. We expect the surge in token consumption on rising demand for AI models and agents to benefit CSPs like **Baidu, Alibaba, Tencent (BAT), Kingsoft Cloud (KC), AI labs and Kling**. According to analysis from Artificial Analysis: (1) there is a narrowing intelligence gap between US and China models; (2) API cost is just a fraction vs US benefiting from MoE, linear attention, and MFU; (3) Coding, Workspace scenario and mid- to long-form content should see significant opportunities ahead

AI Models

According to OpenRouter, token consumption went up significantly in Feb-26 due to (a) a **surge in adoption of coding / agents**, which allows people using AI from "to chat" to "to work". It connects with large models which can be deployed locally or on servers, which allows you to automate workflow through natural language chat; (b) **release of a number of high-performance and cost-efficient models** which include Kimi K2.5 (Jan-26), M2.5 (Feb-26) and GLM-5 (Feb-26).


According to OpenRouter, rising adoption of Coding / Agents and release of high-performance and cost-efficient China models are driving token consumption

Exhibit 6 - Top 5 ranking in terms of token consumption by model the latest week

No	Company	Model	Tokens
1	DeepSeek	DeepSeek V4 Flash	7.22T
2	Xiaomi	MiMo-V2.5	6.3T
3	Tencent	Hy3	4.82T
4	DeepSeek	DeepSeek V4 Pro	3.28T
5	Z.AI	GLM 5.2	2.89T

Source: OpenRouter, Jefferies

Exhibit 7 - Weekly token consumption over the past 12 months as of 3 Aug (tn)




Source: OpenRouter, Jefferies

According to OpenRouter (between 27 Jul and 2 Aug), China models token consumption decreased by 14.8% compared to prior week at 28tn. This compares to US models at 4.4tn token consumption. DeepSeek V4 Flash weekly token consumptions under rapid growth.

In OpenRouter, China models surpassed US models in terms of token consumption the week of 27 Jul


Exhibit 8 - Market share: weekly token consumption by company as of 3 Aug



Source: OpenRouter, Jefferies

On token consumption by company, Google (26.1%) comes first, followed by DeepSeek (24.9%), with the next three accounting for around 27%. The top 5 in weekly token consumption by model are DeepSeek V4 Flash (7.22tn), MiMo V2.5 (6.3tn), and Hy3 (4.82tn), followed by DeepSeek V4 Pro (3.28tn) and GLM 5.2 (2.89tn).

Exhibit 9 - Weekly token consumption since Jan-26 (bn)



Source: OpenRouter, Jefferies

DeepSeek V4 Flash ranked first in terms of token consumption the week of 27 Jul

Exhibit 10 - Top 5 ranking in terms of token consumption by model the latest week

No	Company	Model	Tokens
1	DeepSeek	DeepSeek V4 Flash	7.22T
2	Xiaomi	MiMo-V2.5	6.3T
3	Tencent	Hy3	4.82T
4	DeepSeek	DeepSeek V4 Pro	3.28T
5	Z.AI	GLM 5.2	2.89T

Source: OpenRouter, Jefferies

Exhibit 11 - No 11-20 ranking in terms of token consumption by model in latest week

No	Company	Model	Tokens
11	inclusionAI	Ling-3.0-flash (free)	1.39T
12	DeepSeek	DeepSeek V4 Flash 0731	1.32T
13	Anthropic	Claude Sonnet 5	1.01T
14	Anthropic	Claude Sonnet 4.6	998B
15	Google	Gemini 3 Flash Preview	983B
16	Anthropic	Claude Opus 5	981B
17	Poolside	Laguna S 2.1 (free)	930B
18	Google	Gemini 2.5 Flash Lite	653B
19	Xiaomi	MiMo-V2.5-Pro	595B
20	Anthropic	Claude Opus 4.8	558B

Source: OpenRouter, Jefferies


Based on Artificial Analysis, China models DeepSeek, Qwen, Kimi, MiniMax and Zhipu are closing the gap with US models in terms of intelligence. China models are highly cost-efficient on Mixture

China models are narrowing the gap with US models in terms of intelligence. On Input/

of Experts (MoE) architecture, linear attention, and MLU (Model Flog Utilization). According to 2026 Stanford AI Index Report (Apr-26), the top US model leads Chinese models by 2.7%. Separately, Anthropic (1503), xAI (1495), Google (1494), Alibaba (1449) and DeepSeek (1424) occupy the top tier of Arena Elo ratings as of Mar-26.

Output API, China large models are highly efficient vs US models.


Exhibit 12 - Artificial Intelligence vs Blended Price



Source: Artificial Analysis, Jefferies

Exhibit 13 - Quarterly Input / Output prices for China models


(USD per million tokens)



Source: Artificial Analysis, Jefferies

Exhibit 14 - Quarterly Input / Output prices for US models


(USD per million tokens)



Source: Artificial Analysis, Jefferies

Exhibit 15 - Cost per Intelligence Index Task as of 3 Aug

● Answer ● Reasoning ● Cache Write ● Cache Hit ● Input




Source: Artificial Analysis, Jefferies

Exhibit 16 - Cache hit/Input/Output API price for leading models as of 3 Aug

● Cache Hit ● Input ● Output

Artificial Analysis




Source: Artificial Analysis, Jefferies

In terms of Input/Output API price, China models are just a fraction vs US models.

Exhibit 17 - Intelligence index for top 25 models as of 3 Aug

Artificial Analysis




Among top 20 large models, there are 7 which are China models and open weighted as of 3 Aug.

Exhibit 18 - Frontier Language Model Intelligence over time as of 3 Aug


● Anthropic ● Kimi ● Z AI ● Google ● DeepSeek ● Alibaba ● MiniMax

Artificial Analysis




Source: Artificial Analysis, Jefferies

Exhibit 20 - Coding index for top 25 models as of 3 Aug



Source: Artificial Analysis, Jefferies


Exhibit 22 - Output speed to run Artificial Intelligence Index as of 3 Aug



Source: Artificial Analysis, Jefferies


Model Input/Output Price By Companies

Exhibit 23 - DeepSeek: Input price



Source: Artificial Analysis, Jefferies


Exhibit 25 - Zhipu: Input price



Source: Artificial Analysis, Jefferies


Note: API price is the same across all context lengths in GLM-5.2 vs tiered pricing in GLM-5.1

Exhibit 27 - Qwen: Input price




Source: Artificial Analysis, Jefferies

Exhibit 24 - DeepSeek: Output price



Source: Artificial Analysis, Jefferies


Exhibit 26 - Zhipu: Output price



Source: Artificial Analysis, Jefferies


Note: API price is the same across all context lengths in GLM-5.2 vs tiered pricing in GLM-5.1

Exhibit 28 - Qwen: Output price




Source: Artificial Analysis, Jefferies

Exhibit 29 - MiniMax: Input price




Source: Artificial Analysis, Jefferies
Note: > 512k input tokens for MiniMax-M3

Exhibit 31 - Xiaomi: Input price




Source: Artificial Analysis, Jefferies

Exhibit 33 - Tencent: Input price




Source: Artificial Analysis, Jefferies

Exhibit 35 - Kimi: Input price




Source: Artificial Analysis, Jefferies

Exhibit 30 - MiniMax: Output price




Source: Artificial Analysis, Jefferies
Note: > 512k input tokens for MiniMax-M3

Exhibit 32 - Xiaomi: Output price




Source: Artificial Analysis, Jefferies

Exhibit 34 - Tencent: Output price




Source: Artificial Analysis, Jefferies

Exhibit 36 - Kimi: Output price




Source: Artificial Analysis, Jefferies

Exhibit 37 - Doubao Seed: Input price




Source: Company, Jefferies
Note: 128k < input tokens <= 256k for doubao-seed-2.0-pro, doubao-seed-1.8

Exhibit 39 - Baidu: Input price




Source: Artificial Analysis, Jefferies
Note: 32k < Input tokens <= 128k

Exhibit 41 - StepFun: Input price




Source: Artificial Analysis, Jefferies

Exhibit 38 - Doubao Seed: Output price




Source: Company, Jefferies
Note: 128k < input tokens <= 256k for doubao-seed-2.0-pro, doubao-seed-1.8

Exhibit 40 - Baidu: Output price




Source: Artificial Analysis, Jefferies
Note: 32k < Input tokens <= 128k

Exhibit 42 - StepFun: Output price




Source: Artificial Analysis, Jefferies

Exhibit 43 - Anthropic Claude Opus: Input price




Source: Artificial Analysis, Jefferies

Exhibit 45 - Anthropic Claude Sonnet: Input price




Source: Artificial Analysis, Jefferies

Exhibit 44 - Anthropic Claude Opus: Output price



Source: Artificial Analysis, Jefferies

Exhibit 46 - Anthropic Claude Sonnet: Output price



Source: Artificial Analysis, Jefferies

Model Pricing Trends

Exhibit 47 - Model pricing trends

Model Provider	Model	Launch Date	Contract	Jan-26		Jul-26		Aug-26	
				Input (\$M) (million tokens)	Output (\$M) (million tokens)	Input (\$M) (million tokens)	Output (\$M) (million tokens)	Input (\$M) (million tokens)	Output (\$M) (million tokens)
Anthropic	Qwen3.8-Max	8/3/2026	(5, 1M)						
	Qwen3.7-Max	5/10/2026	(5, 1M)	4 (original 12, promo 18 23 Jul)	18 (original 36, promo 18 23 Jul)	2	18 (original 36, promo 18 23 Jul)	12	36
	Qwen3.7-Plus	5/10/2026	(5, 256K)	1.6 (original 2, promo 18 23 Jul)	64 (original 6, promo 18 23 Jul)	6	6	original 12 (limited 20% off)	original 18 (limited 20% off)
	Qwen3.6-Max Preview	4/29/2026	(5, 128K)	5	54	5	24	original 12 (limited 20% off)	original 24 (limited 20% off)
	Qwen3.6-Max	3/18/2026	(5, 256K)	15	90	10	60	15	90
	Qwen3.5-Plus	2/16/2026	(5, 128K)	8	48	8	48	8	48
Tencent	Hyper	5/6/2026	(55k, 1M)	2	12	2	12	2	12
	HF 3.0 Preview	4/29/2026	(5, 14K)	1.2	4	1.2	4	1.2	4
	HF 2.0 Think	1/16/2025	(5, 13K)	1.9/5	15.9	1.9/5	15.9	1.9/5	15.9
	HF 2.0 Instruct	11/1/2025	(5, 13K)	2.1/6	7.1/6	2.1/6	7.1/6	1.9/6	7.1/6
	Hyper-3T	3/21/2025	(25K, 128K)	1	1	1	4	1	4
	Hyper-3T Turbo	2/16/2025	(25K, 128K)	0.8	8	0.8	8	0.8	8
Baidu	ERNIE 4.5 Turbo	6/2/2026	(5, 32K)	1.2	16	1.2	16	1.2	16
	ERNIE 4.5 Turbo-VL	6/2/2026	(5, 32K)	6	22	6	22	6	22
	ERNIE 4.5 Turbo-VL	6/2/2026	(5, 32K)	6	22	6	22	6	22
	ERNIE 4.5 Turbo-VL	6/2/2026	(5, 32K)	10	40	10	40	10	40
	ERNIE 4.5 Turbo-VL	6/2/2026	(5, 32K)	2	9	2	9	2	9
	ERNIE 4.5 Turbo-VL	6/2/2026	(5, 32K)	1.2	16	1.2	16	1.2	16
ByteDance	Doubao-vend-2.0-pro	2/14/2026	(5, 128K)	4.8	24	4.8	24	4.8	24
	Doubao-vend-2.0-ba	2/14/2026	(5, 128K)	9.6	48	9.6	48	9.6	48
	Doubao-vend-2.0-max	2/14/2026	(5, 128K)	0.9	5.4	0.9	5.4	0.9	5.4
	Doubao-vend-2.0-max	2/14/2026	(5, 128K)	1.8	10.8	1.8	10.8	1.8	10.8
	Doubao-vend-2.0-max	2/14/2026	(5, 128K)	0.2	1	0.2	1	0.2	1
	Doubao-vend-2.0-max	2/14/2026	(5, 128K)	0.8	4	0.8	4	0.8	4
Microsoft	MoMoMo M2	6/1/2026	Input (512x1)	4.2	16.8	4.2	16.8	4.2	16.8
	MoMoMo M2.7	5/10/2026	Input (512x1)	4.2	16.8	4.2	16.8	4.2	16.8
	MoMoMo M2.7	5/10/2026	Input (512x1)	4.2	16.8	4.2	16.8	4.2	16.8
	MoMoMo M2.7	5/10/2026	Input (512x1)	4.2	16.8	4.2	16.8	4.2	16.8
	MoMoMo M2.7	5/10/2026	Input (512x1)	4.2	16.8	4.2	16.8	4.2	16.8
	MoMoMo M2.7	5/10/2026	Input (512x1)	4.2	16.8	4.2	16.8	4.2	16.8
Zhipu AI	GLM-5.1	5/1/2026	(5, 32K)	6	24	6	24	6	24
	GLM-5.1 Turbo	5/1/2026	(5, 32K)	5	20	5	20	5	20
	GLM-5.1	5/1/2026	(5, 32K)	5	20	5	20	5	20
	GLM-5.1	5/1/2026	(5, 32K)	5	20	5	20	5	20
	GLM-5.1	5/1/2026	(5, 32K)	5	20	5	20	5	20
	GLM-5.1	5/1/2026	(5, 32K)	5	20	5	20	5	20
Kimi	Kimi K2	7/16/2026							
	Kimi K2.5	6/10/2026							
	Kimi K2.5	6/10/2026							
	Kimi K2.5	6/10/2026							
	Kimi K2.5	6/10/2026							
	Kimi K2.5	6/10/2026							
DeepSeek	DeepSeek V3	5/28/2026							
	DeepSeek V3	5/28/2026							
	DeepSeek V3	5/28/2026							
	DeepSeek V3	5/28/2026							
	DeepSeek V3	5/28/2026							
	DeepSeek V3	5/28/2026							
Arboreus	Arboreus V1	5/10/2026							
	Arboreus V1	5/10/2026							
	Arboreus V1	5/10/2026							
	Arboreus V1	5/10/2026							
	Arboreus V1	5/10/2026							
	Arboreus V1	5/10/2026							

Source: Company, Jefferies

Video generation model offerings

Red color denotes price hike / new price during the month

Pricing plans for different video generation models

Exhibit 48 - Video generation model pricing in RMB

App	Model	Monthly standard plan (RMB)	Monthly credit quota	Credit per video	Length per video	Cost (RMB) per video
HappyHorse AI	HappyHorse-1.1 (720p)		RMB 0.9/s		3-15s	1.62-8.1
	HappyHorse-1.1 (1080p)		RMB 1.2/s		3-15s	2.16-10.8
	HappyHorse-1.0	70	875	19-95	3-15s	1.52-7.6
Kling AI	Video 3.0/3.0 Omni - 720p	58	660	27-135	3-15s	2.37-11.86
	Video 3.0/3.0 Omni - 1080p	58	660	36-180	3-15s	3.16-15.82
	Video O1 - 720p	58	660	30-60	5-10s	2.64-5.27
	Video O1 - 1080p	58	660	40-80	5-10s	3.52-7.03
	Video 2.6	58	660	50-100	5-10s	4.39-8.79
	Video 2.5 Turbo - 720p	58	660	15-30	5-10s	1.32-2.64
Jimeng AI	Video 2.5 Turbo - 1080p	58	660	25-50	5-10s	2.2-4.39
	Seedance 2.0 Fast	69	725	20-75	4-15s	1.9-7.14
	Seedance 2.0	69	725	32-120	4-15s	3.05-11.42
	Seedance 1.5 Pro	69	725	50-120	5-12s	4.76-11.42
	Seedance 1.0	69	725	50-100	5-10s	4.76-9.52
	Hailuo 2.3 - 768p	68	1000	50-100	6-10s	3.4-6.8
Hailuo AI	Hailuo 2.3 - 1080p	68	1000	160	6s	10.88
	Hailuo 2.3 Fast - 768p	68	1000	30-60	6-10s	2.04-4.08
	Hailuo 2.3 Fast - 1080p	68	1000	100	6s	6.8
	Hailuo 2.0 - 512p	68	1000	24-50	6-10s	1.63-3.4
	Hailuo 2.0 - 768p	68	1000	50-100	6-10s	3.4-6.8
	Hailuo 2.0 - 1080p	68	1000	160	6s	10.88
Vidu	Vidu Q3 Pro - 540p	59	800	10-160	1-16s	0.74-11.8
	Vidu Q3 Pro - 720p	59	800	25-400	1-16s	1.84-29.5
	Vidu Q3 Pro - 1080p	59	800	30-480	1-16s	2.21-35.4
	Vidu Q3 Turbo - 540p	59	800	8-128	1-16s	0.59-9.44
	Vidu Q3 Turbo - 720p	59	800	12-192	1-16s	0.89-14.16
	Vidu Q3 Turbo - 1080p	59	800	14-224	1-16s	1.03-16.52
	Vidu Q3 reference - 540p	59	800	30-160	3-16s	2.21-11.8
	Vidu Q3 reference - 720p	59	800	60-320	3-16s	4.43-23.6
	Vidu Q3 reference - 1080p	59	800	75-400	3-16s	5.53-29.5

Source: Companies, Jefferies

Exhibit 49 - Video generation model pricing in USD

App	Model	Monthly standard plan (US\$)	Monthly credit quota	Credit per video	Length per video	Cost (US\$) per video
Kling AI	Video 3.0/3.0 Omni - 720p	8.8	660	27-135	3-15s	0.36-1.8
	Video 3.0/3.0 Omni - 1080p	8.8	660	36-180	3-15s	0.48-2.4
	Video O1 - 720p	8.8	660	30-60	5-10s	0.4-0.8
	Video O1 - 1080p	8.8	660	40-80	5-10s	0.53-1.07
	Video 2.6	8.8	660	50-100	5-10s	0.67-1.33
	Video 2.5 Turbo - 720p	8.8	660	15-30	5-10s	0.2-0.4
Dreamina AI	Video 2.5 Turbo - 1080p	8.8	660	25-50	5-10s	0.33-0.67
	Seedance 2.0 Fast	15	1575	56-210	4-15s	0.53-2
	Seedance 2.0	15	1575	68-225	4-15s	0.65-2.14
	Seedance 1.5 Pro	15	1575	100-340	5-12s	0.95-2.29
	Seedance 1.0 Mini	15	1575	25-50	5-10s	0.24-0.48
	Hailuo 2.3 - 1080p	9.99	1000	160	6s	1.6
Hailuo AI	Hailuo 2.3 Fast - 768p	9.99	1000	30-60	6-10s	0.3-0.6
	Hailuo 2.3 Fast - 1080p	9.99	1000	100	6s	1
	Hailuo 2.0 - 768p	9.99	1000	50-100	6-10s	0.5-1
	Hailuo 2.0 - 1080p	9.99	1000	160	6s	1.6
Google Gemini	Veo 3.1 Lite	With Google AI Plus (US\$ 7.99 per month with 200 credits) and Pro (US\$ 19.99 per month with 1000 credits) plan			4-8s	
	Veo 3.1	With Google AI Ultra plan at US\$249.99 per month with 25,000 credits			4-8s	
PixVerse	PixVerse V6 - 720p	10	1200	10-150	1-15s	0.08-1.25
	PixVerse V6 - 1080p	10	1200	20-300	1-15s	0.17-2.5
	PixVerse C1 - 720p	10	1200	10-150	1-15s	0.08-1.25
	PixVerse C1 - 1080p	10	1200	20-300	1-15s	0.17-2.5
	PixVerse V5.6 - 720p	10	1200	10-100	1-10s	0.08-0.83
	PixVerse V5.6 - 1080p	10	1200	20-160	1-8s	0.17-1.33
	PixVerse V5 - 720p	10	1200	8-80	1-10s	0.07-0.67
	PixVerse V5 - 1080p	10	1200	16-128	1-8s	0.13-1.07
	Gen 4.5	15	625	24-120	2-10s	0.58-2.88
	Gen 4 Video	15	625	60-120	5-10s	1.44-2.88
Runway	Gen 4 Video Turbo	15	625	25-50	5-10s	0.6-1.2

Source: Companies, Jefferies

Exhibit 50 - Video generation membership comparison (original price)

App	Membership	Monthly price (RMB)	Cost per credit (RMB)	Cost per second (720P)	Cost per second (1080P)
HappyHorse AI	Standard	87.5	0.10	RMB 0.9/s	RMB 1.6/s
	Pro	350	0.10	RMB 0.9/s	RMB 1.6/s
	Gold	66	0.10	RMB 0.9/s	RMB 1.2/s
Kling AI - Video 3.0	Platinum	268	0.09	RMB 0.8/s	RMB 1.07/s
	Diamond	666	0.08	RMB 0.75/s	RMB 1.0/s
	Black Gold	1314	0.05	RMB 0.45/s	RMB 0.61/s
Jimeng AI - Seedance 2.0	Basic	79	0.11		RMB 0.87/s
	Standard	239	0.11		RMB 0.87/s
	Pro	649	0.11		RMB 0.84/s

Source: Company, Jefferies

Exhibit 51 - Video generation membership comparison (after promotion)

App	Membership	Monthly price (RMB)	Cost per credit (RMB)	Cost per second (720P)	Cost per second (1080P)
HappyHorse AI	Standard	70	0.08	RMB 0.51/s	RMB 0.9/s
	Pro	245	0.07	RMB 0.44/s	RMB 0.78/s
	Gold	58	0.09	RMB 0.79/s	RMB 1.05/s
Kling AI - Video 3.0	Platinum	234	0.08	RMB 0.70/s	RMB 0.94/s
	Diamond	586	0.07	RMB 0.66/s	RMB 0.88/s
	Black Gold	1314	0.05	RMB 0.45/s	RMB 0.61/s
Jimeng AI - Seedance 2.0	Basic	69	0.10		RMB 0.76/s
	Standard	199	0.09		RMB 0.72/s
	Pro	499	0.08		RMB 0.65/s


Source: Company, Jefferies

OpenClaw

According to OpenRouter, top 3 models for OpenClaw based on token consumption come from China, namely, MiniMax M3, DeepSeek V4 Flash, and Step 3.7 Flash. Other models include DeepSeek V4 Pro, GLM 5.2, and Hy3.

According to OpenRouter, MiniMax M3, DeepSeek V4 Flash, and Step 3.7 Flash are ranked top 3 in terms of token consumption on OpenClaw

Exhibit 52 - Token consumption for OpenClaw (bn)




Source: OpenRouter, Jefferies

Exhibit 54 - Top 15 models used by OpenClaw over the last 30 days (as of 3 Aug)

No	Model	Tokens
1	MiniMax M3	901B
2	DeepSeek V4 Flash	531B
3	Step 3.7 Flash	328B
4	GLM 5.2	256B
5	DeepSeek V4 Pro	227B
6	Nemotron 3 Super	204B
7	Claude Sonnet 5	141B
8	Hy3	133B
9	Claude Sonnet 4.6	119B
10	Nemotron 3 Ultra	118B
11	Claude Opus 4.8	81.2B
12	Claude Opus 4.6	73.5B
13	Gemini 2.5 Flash Lite	64.1B
14	MiMo-V2.5-Pro	62.1B
15	MiMo-V2.5	56.5B

Source: OpenRouter, Jefferies

Exhibit 53 - Token consumption for Hermes Agent (bn)



Source: OpenRouter, Jefferies


Exhibit 55 - Top 15 models used by Hermes Agent over the last 30 days (as of 3 Aug)

No	Model	Tokens
1	Hy3	5.69T
2	DeepSeek V4 Flash	5.4T
3	Step 3.7 Flash	2.61T
4	DeepSeek V4 Pro	2.48T
5	Nemotron 3 Ultra	2.32T
6	MiniMax M3	2.11T
7	GLM 5.2	1.89T
8	Nemotron 3 Super	708B
9	Claude Sonnet 5	578B
10	Kimi K3	545B
11	MiMo-V2.5-Pro	356B
12	Claude Opus 4.8	329B
13	Ling-3.0-flash (free)	321B
14	MiMo-V2.5	319B
15	Laguna S 2.1	288B

Source: OpenRouter, Jefferies

According to PinchBench, Anthropic's Claude Opus 4.8 Fast ranked first in terms of success rate (93.5%), followed by Qwen3.7 Max (92.5%) and Claude Opus 4.8 (90.5%).


Exhibit 56 - Global: Top 10 models by success rate as of 3 Aug



Source: PinchBench, Jefferies

In Chinese market, Qwen3.7 Max (92.5%) ranked first, followed by MiMo V2.5 (89.7%), Qwen3.6 Flash (88.1%) and MiMo V2.5 Pro (87.5%).

Exhibit 57 - China: Top 10 models by success rate as of 3 Aug




Source: PinchBench, Jefferies

AI Assistants

- MoM trend: DAU for Qwen reached 29.2m in Jun (vs May: 27.9m) and Yuanbao reached 8.6m (vs May: 8.8m). For Doubao, it increased by 6% MoM to 167m (vs May: 158m). For DeepSeek, DAU increased by about 6% MoM to about 32.4m.


DAU for a number of leading players increased MoM in Jun

Exhibit 58 - Weekly DAU for leading AI Assistants in China (mn)




Source: QuestMobile, Jefferies

Exhibit 60 - Weekly user time spent per weekly DAU for leading AI Assistants in China (mins)




Source: QuestMobile, Jefferies

Exhibit 59 - Monthly DAU for Leading AI Assistants in China (mn)



Source: QuestMobile, Jefferies

Exhibit 61 - Monthly user daily time spent per monthly DAU for leading AI Assistants in China (mins)



Source: QuestMobile, Jefferies

Video generation models

According to Artificial Analysis (as of 3 Aug), MiniMax H3 and HappyHorse-1.1 ranked top 5 in terms of text to video and image to video. Others include Gemini Omnni Flash, Seedance 2.0 720p, Wan 2.7, and grok-imagine-video.

Exhibit 62 - Top 15 image to video leaderboard (with audio)

No	Model	ELO
1	Dreamina Seedance 2.0 720p	1,200
2	Gemini Omnni Flash	1,199
3	MiniMax H3	1,118
4	grok-imagine-video-1.5-preview	1,110
5	HappyHorse-1.1	1,100
6	Wan 2.7	1,093
7	HappyHorse-1.0	1,088
8	SkyReels V4	1,085
9	Veo 3.1	1,082
10	grok-imagine-video	1,078
11	Veo 3.1 Fast	1,075
12	Kling 3.0 1080p (Pro)	1,073
13	PixVerse V6	1,070
14	Kling 3.0 720p (Standard)	1,066
15	Vidu Q3 Pro	1,063

Source: Artificial Analysis, Jefferies

Exhibit 63 - Top 15 text to video leaderboard (with audio)

No	Model	ELO
1	Gemini Omnni Flash	1,244
2	MiniMax H3	1,239
3	Dreamina Seedance 2.0 720p	1,224
4	Wan 2.7-260612	1,160
5	HappyHorse-1.1	1,149
6	HappyHorse-1.0	1,127
7	Kling 3.0 1080p (Pro)	1,111
8	Wan 2.7	1,108
9	SkyReels V4	1,108
10	Kling 3.0 720p (Standard)	1,101
11	Veo 3.1	1,094
12	Kling 3.0 Omnni 720p (Standard)	1,092
13	Kling 3.0 Omnni 1080p (Pro)	1,092
14	Veo 3.1 Fast	1,091
15	Veo 3.1 Lite	1,090

Source: Artificial Analysis, Jefferies

Global and China AI Mobile App MAU Ranking

Exhibit 64 - Global AI mobile app MAU ranking (Jun-26)

Rank	App	MAU (m)	MoM
1	ChatGPT	995.49	0.78%
2	Doubao	324.31	-1.71%
3	Qwen	251.13	7.46%
4	Gemini	162.74	7.61%
5	Quark	159.24	-1.78%
6	DeepSeek	139.08	-1.50%
7	Yuanbao	103.92	-7.05%
8	Baidu Drive	101.03	-1.35%
9	Jimeng AI	87.36	-3.33%
10	Dola	79.36	1.99%
11	Meta view	73.06	4.02%
12	Grok	67.88	-7.30%
13	Nova	59.09	-5.10%
14	Claude	39.81	10.79%
15	AI Mirror	34.9	0.05%

Source: AICPB, Jefferies

Exhibit 65 - China AI Mobile App MAU Ranking (Jun-26)

Rank	App	MAU (m)	MoM
1	Doubao	324.31	-1.71%
2	Qwen	251.13	7.46%
3	Quark	159.24	-1.78%
4	DeepSeek	139.08	-1.50%
5	Yuanbao	103.92	-7.05%
6	Baidu Drive	101.03	-1.35%
7	Jimeng AI	87.36	-3.33%
8	Kimi AI	22.69	-12.40%
9	Kling	8.11	-1.69%
10	Nano AI Search	8.08	-4.46%
11	Baidu Wenku	7.39	-5.72%
12	ERNIE Bot	6.57	-6.55%
13	Lovekey Keyboard	4.91	-1.09%
14	Zhipu Qingyan	4.82	-3.93%
15	Cat Box	3.91	-0.30%

Source: AICPB, Jefferies

According to aicpb.com, MAU for China AI apps, which include Doubao, Qwen, Quark, and Yuanbao, experienced strong MoM growth due to Chinese New Year in Feb

In domestic market, MAU growth for Qwen experienced fast MoM growth in Jun


Jun 26: Monthly App Tracker from QuestMobile Analysis

Online shopping

- MoM trend: DAU for Taobao/VIPS/PDD decreased by 2%/7%/3% respectively in Jun. DAU for JD increased by 7% MoM in Jun.
- YoY trend: DAU for JD/Taobao decreased 2%/2% YoY respectively, and PDD increased by 3% YoY.


PDD maintained fast DAU growth on YoY basis in Jun

Exhibit 66 - MAU for Taobao / PDD / JD in Jun-26



Source: QuestMobile, Jefferies

Exhibit 67 - DAU for Taobao / PDD / JD in Jun-26




Source: QuestMobile, Jefferies

Food delivery

- MoM trend: Meituan MAU maintained stable during the month.
- YoY trend: Meituan MAU increased by 2% (vs 3% YoY in May) in Jun.

Meituan MAU was largely stable MoM in Jun

Exhibit 68 - MAU for Meituan app in past 12 months



Source: QuestMobile, Jefferies


Online Music

- MoM trend: MAU for Soda Music remained stable MoM to 167.4m in Jun. For TME, MAU for QQ Music/Kugou/Kuwo decreased by 2%/4%/3% MoM respectively in Jun. MAU for NetEase Cloud Music decreased 1% MoM during the month.

Soda Music maintains fast MAU YoY growth while largely stable MoM in Jun

- YoY trend: MAU for Soda Music increased by 69% YoY while NetEase Cloud Music MAU increased by 1% YoY in Jun. For TME, MAU for QQ Music/Kugou/Kuwo decreased by 1%/11%/18% YoY respectively.

Exhibit 69 - MAU for TME, Soda Music and NetEase Cloud Music in Jun-26




Source: QuestMobile, Jefferies

Online video

- MoM trend: MAU for Tencent Video/iQ/Mango TV/Youku decreased by 2%/2%/2%/4% MoM.
- YoY trend: MAU for leading long form video players declined YoY during the month.

Exhibit 71 - MAU for IQ, Tencent Video and Youku in Jun-26




Source: QuestMobile, Jefferies

Live-streaming and short video

- MoM trend: DAU for Douyin main app increased 1% MoM in Jun. For game broadcasting, MAU for Huya decreased by 4% and Douyu decreased by 1% MoM in Jun. For BILI, MAU increased by 1% MoM during the month.
- YoY trend: DAU for Douyin main app maintains solid growth at 20% YoY. For game broadcasting, both Huya and Douyu experienced 8% and 9% YoY decline in MAU respectively. For BILI, MAU increased by about 4% YoY during the month.

Exhibit 70 - MAU for Soda Music in past 12 months




Source: QuestMobile, Jefferies

MAU for leading long form video platform declined YoY in Jun


Douyin DAU maintained solid YoY growth in Jun

Exhibit 72 - DAU for Douyin over the past 12 months



Source: QuestMobile, Jefferies

Exhibit 73 - MAU for BILI over the past 12 months




Source: QuestMobile, Jefferies

Online travel

- MoM trend: MAU for TCOM / Qunar declined by 1%/1% MoM respectively in Jun.
- YoY trend: MAU for TCOM/Qunar declined by 8%/9% YoY in Jun.

TCOM MAU in domestic market decreased MoM in Jun after CNY in Feb

Exhibit 74 - TCOM's MAU over the past 12 months



Source: QuestMobile, Jefferies

News app

- MoM: MAU for Weibo remained stable and MAU for Toutiao decreased by 2% in Jun.
- YoY: MAU for Weibo remained stable YoY while Toutiao decreased by 6% in Jun.

MAU for Weibo maintained stable YoY growth in Jun

Exhibit 75 - Weibo MAU over the past 12 months




Source: QuestMobile, Jefferies

Baidu app

- MoM: MAU declined by 1% MoM and DAU decreased by 1% in Jun.
- YoY: MAU declined by 8% YoY and DAU declined by 21% YoY in Jun.

Baidu app MAU experienced declined MoM and YoY in Jun


Exhibit 76 - Baidu MAU over the past 12 months



Source: QuestMobile, Jefferies


IDC estimates on token consumption on public cloud and outlook

Exhibit 77 - China's MaaS market monthly average daily token consumption in 2025




Source: IDC, Jefferies

Exhibit 79 - Market share of major vendors in China's MaaS market in 2025 (revenue basis)




Source: IDC, Jefferies

Exhibit 81 - The most important factors that drive people's decision to deploy model widely




Source: IDC, Jefferies

Exhibit 78 - China's MaaS market share by token consumption in 2025




Source: IDC, Jefferies

Exhibit 80 - Market share of major vendors in China's private large model platform market in 2025 (revenue basis)



Source: IDC, Jefferies

Exhibit 82 - China's MaaS market scale and forecast 2024-2030 (token basis)



Source: IDC, Jefferies

Appendix

Omdia heatmap for agentic AI development platform in Asia and Oceania

According to Omdia, agentic AI software market is expected to increase at 105% CAGR from USD271m in 2025 to USD9.7bn by 2030. Alibaba Cloud captures the highest rankings in 5 areas: (1) context engineering, (2) model support, (3) multi-agent framework, (4) operation and lifecycle management, and (5) open source and community.

According to Omdia, Alibaba Cloud is ranked as one of the leaders and top choice for building and deploying AI agents. It captures the highest ranking for 5 out of 7 metrics.

Exhibit 83 - Omdia heatmap for agentic AI development platform in Asia and Oceania

	LEADER				
	Alibaba Cloud	AWS	Google Cloud Platform	Microsoft Azure	Tencent Cloud
Context engineering					
Model support					
Multi-agent framework					
Operation and lifecycle management					
Tools and marketplace					
Vertical specific solution					
Open source and community					


Source: Omdia, Jefferies

Market size: Coding, Workspace, and Content Creation

Global AI code tools market is expected to be 3x the scale in 2030 vs 2025


According to Grand View Research, the global AI code tools market is expected to increase by 3x to USD26bn from 2025 to 2030. On the other hand, F&S predicts that the China AI code generation market will reach about USD4.7bn in 2028. KA customized solutions are expected to be a key revenue contributor.

Exhibit 84 - Global AI code tools market size (USDbn)




Source: Grand View Research, Jefferies

Exhibit 85 - China AI code generation market size by service type (RMBbn)



Source: Frost & Sullivan, Jefferies

Exhibit 86 - China AI code penetration in different industries




In terms of AI code penetration, internet, gaming, AI, finance, and energy are top 5 key sectors

Source: Frost & Sullivan, Jefferies

According to Gonyon Industry Research (2024), market size of AI+Workspace software is expected to reach RMB191bn in 2028, 6x the scale vs 2024.


Exhibit 87 - China AI + Workspace Market Size



Source: Gonyon, Jefferies


According to Precedence Research, global generative AI content creation market is expected to increase by 7x surpassing USD140bn in 2035 vs 2025.

Exhibit 88 - Generative AI in Content Creation Market Size 2025 to 2035 (USD bn)



Source: Precedence Research, Jefferies

Exhibit 89 - Daily token consumption in China by enterprises in 2H25 vs 1H25 (tn)




Source: Frost & Sullivan, Jefferies

According to F&S, enterprise daily token consumption increased by 263% HoH to 37tn, among which market share for Alibaba Cloud increased the most and captures no 1 position reaching 32% in 2H25 (1H25:17.7%), followed by 21.3% for Doubao (1H25: 14.1%) and 18.4% for DeepSeek (1H25: 10.3%).


Alibaba's Qwen ranked first in terms of token consumption among enterprises in 2H25

Exhibit 90 - Market share in terms of tokens used among enterprises in 1H25



Source: Frost & Sullivan, Jefferies

Exhibit 91 - Market share in terms of tokens used among enterprises in 2H25




Source: Frost & Sullivan, Jefferies

China Large Language Models

According to F&S / CAICT, China LLM market is expected to increase by about 64% CAGR between 2024 and 2030 and surpass RMB100bn in 2030. By segment, consumer revenue is expected to grow at a slower rate vs enterprise as willingness to pay remains low.


China large language models are expected to increase by 64% CAGR between 2024 to 2030. By segment, on-premise is key revenue contributor

Exhibit 92 - Market size of China LLM (RMB billions)



Source: Frost & Sullivan, CAICT, Jefferies

Exhibit 93 - Market size of China Enterprise LLM (RMB billions)




Source: Frost & Sullivan, CAICT, Jefferies

Anthropic

- In May-26, Anthropic announced a USD965bn post-money valuation on 28 May. It raised USD65bn in Series H funding. ARR crossed USD47bn in early May.
-

Exhibit 94 - ARR of Anthropic



Source: The Information, company reports, Jefferies

Models of Leading Players in China

Exhibit 95 - Alibaba model list

Release Date	Model	Core Functions
Jan-24	Qwen-VL-Plus / Max	Enhanced visual reasoning, high-resolution OCR, and detailed image understanding via API.
Feb-24	Qwen1.5	First major open-source iteration with improved multilingual support and 0.5B to 72B sizes.
Jun-24	Qwen2	Optimized for coding, mathematics, and long-context processing (up to 128K tokens).
Aug-24	Qwen2-Audio	Native audio-to-text and audio-to-audio dialogue capabilities.
Aug-24	Qwen2-VL	Introduced dynamic resolution for visual inputs and long video understanding.
Sep-24	Qwen2.5	Significant upgrades in instruction following, structured data (tables) processing, and coding.
Nov-24	QwQ-32B-Preview	Reasoning-focused model using Reinforcement Learning for complex logic and chain-of-thought.
Jan-25	Qwen2.5-Max	Large-scale MoE (Mixture of Experts) flagship; matches or exceeds GPT-4o benchmarks.
Jan-25	Qwen2.5-VL	"Agentic" vision model capable of identifying and interacting with mobile/PC UI screens.
Mar-25	Qwen2.5-Omni	End-to-end full-modality model supporting seamless real-time text, audio, and visual interaction.
Apr-25	Qwen3 (Series)	New reasoning architecture with hybrid inference modes to reduce hallucinations in complex tasks.
Sep-25	Qwen3-Max-Preview	Trillion-parameter preview version optimized for massive enterprise-scale applications.
Dec-25	Wan2.6-Image	Supports image editing and mixed image-text output.
Jan-26	Qwen3-Max	Frontier reasoning, 1T+ params, agentic workflow mastery.
Feb-26	Qwen3.5	Latest flagship series featuring a 397B MoE model with superior cross-lingual reasoning.
Feb-26	Wan2.6-r2v-flash	Generates multi-shot videos with automatic dubbing, using a character's appearance from a reference video or image.
Feb-26	Qwen3.5-Plus	Outstanding performance comparable to the very latest state-of-the-art models with significant leaps in both pure-text and multimodal capabilities compared to the 3 series.
Feb-26	Qwen3-Coder-Next	Multi-turn tool interactions, future-ready development support.
Mar-26	Qwen3.5-Omni	Support for 256k long-text input and can process more than 10 hours of audio input and over 400 seconds of 720P audio-visual input at 1 FPS.
Apr-26	Wan2.7-Image	The Wan 2.7 image generation and editing model supports text-to-image, text-to-image-set, image-to-image-set, image editing, multi-image reference generation, and interactive editing, with improved performance in text rendering, subject consistency, and complex instruction following. The Pro series supports 4K output; the accelerated version balances quality and response speed.
Apr-26	Qwen3.6-Plus	Native multimodal, 1M context, agentic coding
Apr-26	Happy Oyster	AI model that creates interactive 3D scenes with realistic physics for game, film, and TV.
Apr-26	Qwen3.6-Max-Preview	It brings stronger world knowledge and instruction following, along with significant agentic coding improvements across a wide range of benchmarks.
Apr-26	Qwen-Image-2.0	Delivers a significant leap in visual quality — especially texture detail, lighting, and materials. It supports multilingual in-image text generation and delivers more balanced artistic style performance.
Apr-26	HappyHorse-1.0	It supports all four video generation modalities: Text to Video and Image to Video, each with and without native audio.
May-26	Qwen3.7-Max	Qwen3.7-Max is built to be a versatile agent foundation — equally capable of writing and debugging code, automating office workflows, and sustaining autonomous execution across hundreds or thousands of steps.
Jun-26	Qwen3.7-Plus	A multimodal agent model that unifies vision and language into a single, versatile agent foundation. Building on Qwen3.7's strong text backbone, Qwen3.7-Plus delivers a comprehensive upgrade in vision-language capabilities while retaining full agentic strength in coding, tool use, and productivity workflows.
Jun-26	HappyHorse-1.1	HappyHorse 1.1 capabilities are available via API, providing enterprise customers and developers with a complete integration solution. This release delivers production-ready video synthesis systematically optimized across core content generation scenarios.
Jul-26	Qwen-Image-3.0	Qwen-Image-3.0 supports up to 4.5k token input, effortlessly generating complex layouts such as newspapers, storyboards, and exam papers. It supports precise rendering of text as small as 10px, vividly reproducing details like pores and hair strands with lifelike, micro-level depiction. The model supports native rendering of 12 languages, simulates mainstream interfaces such as web pages, games, and livestreams, and draws on rich world knowledge.
Aug-26	Qwen3.8-Max	Built upon the architectural foundation of Qwen 3.5, Qwen 3.8-Max scales to 2.4 trillion parameters, delivering comprehensive improvements across coding, work, research, and long-horizon tasks. It can not only answer more challenging questions, but also complete complex tasks end-to-end with greater reliability, producing dependable deliverables.

Source: Company, Jefferies

Exhibit 96 - Tencent model list

Release Date	Model	Core Functions
May-24	Hunyuan-DiT	The industry's first Chinese-native Diffusion Transformer (DiT) for high-quality text-to-image generation.
Sep-24	Hunyuan-Turbo	An upgraded LLM based on the Mixture-of-Experts (MoE) architecture, doubling training efficiency and reducing inference costs.
Nov-24	Hunyuan3D 1.0	First open-source 3D generation model supporting text-to-3D and image-to-3D workflows.
Dec-24	HunyuanVideo	A large-scale video generation framework capable of producing high-definition videos with complex motion.
Jan-25	Hunyuan3D 2.0	Significant upgrades in geometry accuracy and texture representation for 3D asset creation.
Feb-25	Hunyuan Turbo S	A "fast-thinking" model using Hybrid-Mamba-Transformer architecture, optimized for near-instant responses and competitive reasoning.
Mar-25	HunyuanVideo-1.2V	An specialized image-to-video model built on the foundation of HunyuanVideo.
Jun-25	Hunyuan3D 2.1	Production-ready 3D asset generation with a new PBR (Physically Based Rendering) material model.
Sep-25	Hunyuan Image 3.0	Flagship image generation model (80B parameters) topping global rankings for semantic understanding.
Nov-25	HunyuanVideo 1.5	A highly efficient and powerful update to the foundation video generation model.
Dec-25	Hunyuan 2.0	Next-generation LLM featuring 406B parameters and a 256K context window for complex reasoning tasks.
Jan-26	Hunyuan Image 3.0-Instruct	An image-to-image editing model supporting precise single-sentence editing and multi-image fusion.
Mar-26	HY-WU	A scalable framework for on-the-fly conditional generation of low-rank(LoRA) updates.
Apr-26	HY-World 2.0	A multi-modal world model framework for world generation and world reconstruction. It accepts diverse input modalities — text, single-view images, multi-view images, and videos — and produces 3D world representations
Apr-26	Hy3 preview	A 295B-parameter Mixture-of-Experts (MoE) model with 21B active parameters and 3.8B MTP layer parameters. It improves significantly on complex reasoning, instruction following, context learning, coding, and agent tasks.
May-26	Hy-MT2	A powerful multilingual model supporting seamless translation across 33 languages
Jul-26	Hy3	Hy3 is a 295B-parameter Mixture-of-Experts (MoE) model with 21B active parameters and 3.8B MTP layer parameters. It significantly outperforms similar-size models and rivals flagship open-source models with 2-5x the parameters. It also shows solid gains in utility across productivity tasks and real-world applications.

Source: Company, Jefferies

Exhibit 97 - Baidu model list

Release Date	Model	Core Functions
Mar-24	ERNIE Speed / Lite / Tiny	Lightweight models optimized for low latency and high cost-efficiency.
Jun-24	ERNIE 4.0 Turbo	Efficiency upgrade for the 4.0 architecture with faster response speeds.
Mar-25	ERNIE 4.5	Baidu's first native multimodal foundation model (text, image, audio, video).
Mar-25	ERNIE X1	Deep-thinking reasoning model designed for complex logic and multi-step tasks.
Apr-25	ERNIE 4.5 Turbo / X1 Turbo	Performance-optimized version of X1 with significantly lower inference costs.
Sep-25	ERNIE X1.1	Major upgrade in factuality, instruction following, and agentic capabilities.
Nov-25	ERNIE 5.0 (Preview)	Preview of next-gen omni-modal model with advanced real-time reasoning.
Jan-26	ERNIE 5.0 (Stable)	Current flagship 2.4T parameter model; top-ranked Chinese model on LMArena.
Apr-26	ERNIE-image	A text-to-image generation model built on a single-stream Diffusion Transformer (DiT) with 8B DiT parameters, achieving leading performance among open weights models on several challenging benchmarks.
May-26	ERNIE 5.1	Significantly reduces pretraining cost while compressing total parameters to ~1/3 and activated parameters to ~1/2 — using only ~6% of the pretraining cost compared to models at similar scale, while achieving leading performance in its class.

Source: Company, Jefferies

Exhibit 98 - MiniMax model list

Release Date	Model	Core Function
Jan-25	MiniMax-Text-01/MiniMax-VL-01	New generation of text and vision language model
Feb-25	T2V-01-Director/I2V-01-Director	Director-level video generation model, better adherence to camera movement description instructions
Apr-25	Image-01	Support generating images of various sizes from text descriptions
Jun-25	MiniMax-M1	80K mind chain and 1M input
Jun-25	MiniMax-Hailuo 02	Native 1080p, SOTA instruction following, extreme physics mastery
Jun-25	Music-1.5 (Beta)	Support music generation by inputting musical inspiration and lyrics
Aug-25	Sppeech-2.5	Significantly enhanced multilingual performance, more realistic and accurate voice cloning, over 40 languages
Sep-25	Music-1.5	Extending song generation to a full 4minutes, one person is the band
Oct-25	MiniMax M2	Open-source MoE, strong code & agent
Oct-25	MiniMax-Hailuo-2.3	Complex video performance & media agent
Oct-25	MiniMax-Speech-2.6	Rea-time response, intelligent parsing, fluent LoRA voice
Oct-25	Music-2.0	A lead forward in model's understanding and expression of music
Dec-25	MiniMax M2.1	Exceptional multilingual capabilities to elevate code engineering
Jan-26	Music-2.5	Creation more precise and music more real
Jan-26	MiniMax M2-Her	Multi-character roleplay, immersive long-horizon interaction
Feb-26	q	Industry-leading coding and reasoning capabilities at a fraction of the cost
Mar-26	Music-2.5+	Enhanced creation experience & efficiency, smarter melodies, purer sound quality
Mar-26	MiniMax M2.7	M2.7 is capable of building complex agent harnesses and completing highly elaborate productivity tasks, leveraging capabilities such as Agent Teams, complex Skills, and dynamic tool search.
Apr-26	Music-2.6	A text-to-audio music generation model, the output supports up to 256kbps bitrate and 44.1kHz sample rate.
Jun-26	MiniMax M3	M3 reaches frontier-level performance on specialized tasks such as coding and agentic work. It uses MSA (MiniMax Sparse Attention), a new attention architecture proposed by our team, and supports ultra-long context windows of up to 1M tokens.
Jul-26	MiniMax H3	MiniMax H3 is an open, general-purpose multimodal video model. It understands text, image, video, and audio inputs in a unified way, and supports video generation, reference-based creation, and video editing.

Source: Company, Jefferies

Exhibit 99 - Moonshot AI model list

Release Date	Model	Core Function
Jan-25	Kimi k1.5	An o1-level multi-modal model
Apr-25	Kimi-VL	Offers advanced multimodal reasoning, long-context understanding, and strong agent capabilities
Apr-25	Kimira-Prover Preview	Can reason in a human-like way and prove mathematical theorems in the Lean 4 language
Apr-25	Kimi-Audio	Open-source audio foundation model excelling in audio understanding, generation, and conversation
Jun-25	Kimi-Dev	A strong and open-source coding LLM for issue resolution
Jun-25	Kimi-Researcher	End-to-end RL training for emerging agentic capabilities
Jul-25	Kimi K2	A state-of-the-art MoE language mode, featuring 32 billion activated parameters and a total of 1 trillion parameters
Sep-25	Kimi-K2-Instruct-0905	The post-trained model best for drop-in, general-purpose chat and agentic experience
Nov-25	Kimi K2 Thinking	Execute up to 200-300 sequential tool calls without human interference, reasoning coherently across hundreds of stpes to solve complex problems
Jan-26	Kimi K2.5	A multimodal model, state-of-the-art coding, vision capabilities, self-directed agent swarm paradigm
Apr-26	Kimi K2.6	Kimi's latest and most intelligent model, possessing stronger and more stable long-term code writing capabilities, significantly improved instruction compliance and self-correction capabilities, and the ability to handle more complex software engineering tasks
Jun-26	Kimi K2.7 Code	Kimi K2.7 Code is a coding-focused agentic model built upon Kimi K2.6. With substantial improvements on real-world long-horizon coding tasks, it strengthens end-to-end task completion across complex software engineering workflows while improving token efficiency, reducing thinking-token usage by approximately 30% compared with Kimi K2.6
Jul-26	Kimi K3	Kimi K3 is a 2.8T-parameter model built on Kimi Delta Attention and Attention Residuals, with native vision capabilities and a 1-million-token context window. It is the world's first open 3T-class model, designed for frontier intelligence across long-horizon coding, knowledge work, and reasoning.

Source: Company, Jefferies

Exhibit 100 - Z.ai model list

Release Date	Model	Core Function
Jul-25	GLM-4.5-Air	Adopts a more streamlined design with 106B total parameters and 12B active parameters.
Jul-25	GLM-4.5	Total parameter count of 355B with 32B active parameters per forward pass
Aug-25	GLM-4.5V	Achieved SOTA performance among open-source VLMs of the same level with 106B parameters and 12B activation parameters
Sep-25	GLM-4.6	Advanced agentic, reasoning and coding capabilities
Dec-25	AutoGLM	Autonomous foundation agents for GUIs
Dec-25	GLM-4.6V	Open source multimodal models with native tool use
Dec-25	GLM-ASR-Nano	A robust, open-source speech recognition model with 1.5B parameters
Dec-25	GLM-TTS	Zero-shot TTS with multi-reward reinforcement learning
Dec-25	GLM-4.7	Complex systems engineering and long-horizon agentic tasks
Jan-26	GLM-4.7-Flash	Open source and free
Feb-26	GLM-OCR	SOTA performance, mastering complex document recognition
Feb-26	GLM-5	MoE flagship, complex engineering & long-horizon agent
Mar-26	GLM-5-Turbo	GLM-5-Turbo is a foundation model deeply optimized for the OpenClaw scenario. It has been specifically optimized for the core requirements of OpenClaw tasks since the training phase, enhancing key capabilities such as tool invocation, command following, timed and persistent tasks, and long-chain execution.
Apr-26	GLM-5V-Turbo	It strengthens multimodal tool use and GUI task execution, allowing the model to better interpret interfaces, design mocks, documents, and charts, plan the next action, and complete end-to-end Agent workflows while maintaining strong pure-text coding and reasoning performance.
Apr-26	GLM-5.1	Built for long-horizon tasks, it runs autonomously for 8 hours, refining strategies through thousands of iterations, and it delivers significant leaps in coding and agentic performance with 6x performance boost. In high-performance database optimization, it reached 21.5k QPS over 600+ iterations and 6,000+ tool calls, which is 6x the performance of a standard 50-turn session
Jun-26	GLM-5.2	GLM-5.2 is Zhipu's most powerful open-source model to date, supporting a truly usable 1M context and maintaining its lead in long-range tasks. It also remains the strongest domestic coding model

Source: Company, Jefferies


Exhibit 101 - Benchmark glossary

Category	Benchmark	Measure	Definition
Knowledge & Reasoning	MMLU-Pro	General knowledge	Measuring Massive Multi-task Language Understanding Professional
	SimpleQA-Verified	Fast accuracy	Simple Question Answering for factuality evaluation
	Chinese-SimpleQA	Chinese facts	Chinese-specific Simple Question Answering benchmark
	GPQA Diamond	Advanced science	Graduate-Level Google-Proof Question Answering
	HLE	Hard reasoning	Humanity's Last Exam for expert-level reasoning
	LiveCodeBench	Live coding	Contamination-free Live Code Benchmarking
	Codeforces	Algorithm solving	Benchmarking Competition-level Code Generation of LLMs
	HMMT 2026 Feb	Math Competition	Feb 2026 Harvard-MIT Mathematics Tournament
	IMOAnswerBench	Olympiad Math	International Mathematical Olympiad Answer-based Benchmark
	Apex	Advanced logic	The AI Productivity Index, a benchmark for assessing whether frontier AI models can perform knowledge work with high economic value
Long Context	Apex Shortlist	Advanced logic	The MathArena Apex Shortlist
	MRCR 1M	Long retrieval	Multi-Round Coreference Resolution at 1 Million tokens
	CorpusQA 1M	Long document QA	A long-context question answering benchmark designed to evaluate models at approximately 1 million token contexts
Agentic	Terminal Bench 2.0	Terminal tasks	A command-line benchmark suite designed to evaluate AI agents on 89 real-world multi-step terminal tasks
	SWE Verified	Code repair	A human-filtered subset of 500 instances from SWE-bench
	SWE Pro	Code repair	A substantially more challenging benchmark that builds upon the best practices of SWE-Bench
	SWE Multilingual	Code repair	A benchmark designed to evaluate the software engineering capabilities of LLMs across multiple programming languages
	BrowseComp	Web browsing	A simple and challenging benchmark that measures the ability of AI agents to locate hard-to-find information
	HLE w/tools	Tool use	Humanity's Last Exam with external tools
	GDPval-AA	Task Planning	A benchmark designed to evaluate AI models on real-world, economically valuable tasks across 44 occupations
	MCPAtlas Public	Multi-tool use	A large-scale benchmark for tool-use competency with real MCP servers
	Toolathlon	Tool Proficiency	A comprehensive benchmark for evaluating language agents on diverse, realistic, and long-horizon tasks

Source: Jefferies


Different third-party analysis shows China is making a meaningful contribution to global token consumption. According to Hugging Face, China models accounted for 41% of global downloads in the 12 months ended Feb-26.

Exhibit 102 - China Daily token usage




Source: National Data Administration (NDA), Jefferies

Exhibit 103 - Doubao daily token consumption



Source: VolcanoEngine, Jefferies

Exhibit 104 - Token share by region



Source: OpenRouter, Jefferies

We would like to thank Fiona Fan, employee of Evalueserve Inc., for providing research support services to our preparation of this report.
We would like to thank Han Wang, employee of Evalueserve Inc., for providing research support services to our preparation of this report.
We would like to thank Erica Qiu, employee of Evalueserve Inc., for providing research support services to our preparation of this report.

Company Valuation/Risks

For Important Disclosure information on companies recommended in this report, please visit our website at <https://javatar.bluematrix.com/sellside/Disclosures.action> or call 212.284.2300.

Analyst Certification:

I, Thomas Chong, certify that all of the views expressed in this research report accurately reflect my personal views about the subject security(ies) and subject company(ies). I also certify that no part of my compensation was, is, or will be, directly or indirectly, related to the specific recommendations or views expressed in this research report.

I, Zoey Zong, certify that all of the views expressed in this research report accurately reflect my personal views about the subject security(ies) and subject company(ies). I also certify that no part of my compensation was, is, or will be, directly or indirectly, related to the specific recommendations or views expressed in this research report.

Registration of non-US analysts: Thomas Chong is employed by Jefferies Hong Kong Limited, a non-US affiliate of Jefferies LLC and is not registered/qualified as a research analyst with FINRA. This analyst(s) may not be an associated person of Jefferies LLC, a FINRA member firm, and therefore may not be subject to the FINRA Rule 2241 and restrictions on communications with a subject company, public appearances and trading securities held by a research analyst.

Registration of non-US analysts: Zoey Zong is employed by Jefferies Hong Kong Limited, a non-US affiliate of Jefferies LLC and is not registered/qualified as a research analyst with FINRA. This analyst(s) may not be an associated person of Jefferies LLC, a FINRA member firm, and therefore may not be subject to the FINRA Rule 2241 and restrictions on communications with a subject company, public appearances and trading securities held by a research analyst. As is the case with all Jefferies employees, the analyst(s) responsible for the coverage of the financial instruments discussed in this report receives compensation based in part on the overall performance of the firm, including investment banking income. We seek to update our research as appropriate, but various regulations may prevent us from doing so. Aside from certain industry reports published on a periodic basis, the large majority of reports are published at irregular intervals as appropriate in the analyst's judgement.

Investment Recommendation Record

(Article 3(1)e and Article 7 of MAR)

Recommendation Published	August 3, 2026 2:32 A.M.
Recommendation Distributed	August 3, 2026 2:32 A.M.

Company Specific Disclosures

James Heaney has a long position in Facebook.

Jefferies Financial Group Inc., its affiliates or subsidiaries make a market in the securities or ADRs of NetEase Cloud Music Inc.

Jefferies Financial Group Inc., its affiliates or subsidiaries make a market in the securities or ADRs of NetEase Inc.

Within the past twelve months, Kuaishou Technology has been a client of Jefferies Financial Group Inc., its affiliates or subsidiaries and investment banking services are being or have been provided.

Jefferies Financial Group Inc., its affiliates or subsidiaries is acting as a manager or co-manager in the underwriting or placement of securities for Kuaishou Technology or one of its affiliates.

Within the past 12 months, Jefferies Financial Group Inc., its affiliates or subsidiaries has received compensation from investment banking services from Kuaishou Technology.

Jefferies International Ltd, its affiliates or subsidiaries has, or had, within the past 12 months an agreement to provide investment services to Kuaishou Technology.

For Important Disclosure information on companies recommended in this report, please visit our website at <https://javatar.bluematrix.com/sellside/Disclosures.action> or call 212.284.2300.

Explanation of Jefferies Ratings

Buy - Describes securities that we expect to provide a total return (price appreciation plus yield) of 15% or more within a 12-month period.

Hold - Describes securities that we expect to provide a total return (price appreciation plus yield) of plus 15% or minus 10% within a 12-month period.

Underperform - Describes securities that we expect to provide a total return (price appreciation plus yield) of minus 10% or less within a 12-month period.

The expected total return (price appreciation plus yield) for Buy rated securities with an average security price consistently below \$10 is 20% or more within a 12-month period as these companies are typically more volatile than the overall stock market. For Hold rated securities with an average security price consistently below \$10, the expected total return (price appreciation plus yield) is plus or minus 20% within a 12-month period. For Underperform rated securities with an average security price consistently below \$10, the expected total return (price appreciation plus yield) is minus 20% or less within a 12-month period.

NR - The investment rating and price target have been temporarily suspended. Such suspensions are in compliance with applicable regulations and/or Jefferies policies.

CS - Coverage Suspended. Jefferies has suspended coverage of this company.

NC - Not covered. Jefferies does not cover this company.

Restricted - Describes issuers where, in conjunction with Jefferies engagement in certain transactions, company policy or applicable securities regulations prohibit certain types of communications, including investment recommendations.

Monitor - Describes securities whose company fundamentals and financials are being monitored, and for which no financial projections or opinions on the investment merits of the company are provided.

Valuation Methodology

Jefferies' methodology for assigning ratings may include the following: market capitalization, maturity, growth/value, volatility and expected total return over the next 12 months. The price targets are based on several methodologies, which may include, but are not restricted to, analyses of market risk, growth rate, revenue stream, discounted cash flow (DCF), EBITDA, EPS, cash flow (CF), free cash flow (FCF), EV/EBITDA, P/E, PE/growth, P/CF, P/FCF, premium (discount)/average group EV/EBITDA, premium (discount)/average group P/E, sum of the parts, net asset value, dividend returns, and return on equity (ROE) over the next 12 months.

Jefferies Franchise Picks

Jefferies Franchise Picks include stock selections from among the best stock ideas from our equity analysts over a 12 month period. Stock selection is based on fundamental analysis and may take into account other factors such as analyst conviction, differentiated analysis, a favorable risk/reward ratio and investment themes that Jefferies analysts are recommending. Jefferies Franchise Picks will include only Buy rated stocks and the number can vary depending on analyst recommendations for inclusion. Stocks will be added as new opportunities arise and removed when the reason for inclusion changes, the stock has met its desired return, if it is no longer rated Buy and/or if it triggers a stop loss. Stocks having 120 day volatility in the bottom quartile of S&P stocks will continue to have a 15% stop loss, and the remainder will have a 20% stop. Franchise Picks are not intended to represent a recommended portfolio of stocks and is not sector based, but we may note where we believe a Pick falls within an investment style such as growth or value.

Risks which may impede the achievement of our Price Target

This report was prepared for general circulation and does not provide investment recommendations specific to individual investors. As such, the financial instruments discussed in this report may not be suitable for all investors and investors must make their own investment decisions based upon their specific investment objectives and financial situation utilizing their own financial advisors as they deem necessary. Past performance of the financial instruments recommended in this report should not be taken as an indication or guarantee of future results. The price, value of, and income from, any of the financial instruments mentioned in this report can rise as well as fall and may be affected by changes in economic, financial and political factors. If a financial instrument is denominated in a currency other than the investor's home currency, a change in exchange rates may adversely affect the price of, value of, or income derived from the financial instrument described in this report. To the extent prices are shown in non-US currency, please note that our local currency price targets are based on a currency conversion using an exchange rate as of the prior trading day (unless otherwise noted). Should there be fluctuations in the exchange rate after this date, that will affect the non-US target prices and should no longer be relied upon. In addition, investors in securities such as ADRs, whose values are affected by the currency of the underlying security, effectively assume currency risk.

Other Companies Mentioned in This Report

- Alibaba Group Holding Limited (9988 HK: HK\$117.00, BUY)
- Alibaba Group Holding Limited (BABA: \$122.25, BUY)
- Alphabet, Inc. (GOOGL: \$356.13, BUY)
- Baidu Inc. (BIDU: \$111.11, BUY)
- Baidu Inc. (9888 HK: HK\$105.90, BUY)
- Bilibili Inc (BILI: \$19.10, BUY)
- Bilibili Inc (9626 HK: HK\$147.70, BUY)
- Kingsoft Cloud (KC: \$11.15, BUY)
- Kuaishou Technology (1024 HK: HK\$43.70, BUY)
- Meituan (3690 HK: HK\$92.50, BUY)
- Meta Platforms, Inc. (META: \$556.71, BUY)
- MiniMax Group Inc. (100 HK: HK\$230.60, BUY)
- NVIDIA Corporation (NVDA: \$200.75, BUY)
- NetEase Cloud Music Inc. (9899 HK: HK\$127.40, BUY)
- NetEase Inc. (NTES: \$133.27, BUY)
- NetEase Inc. (9999 HK: HK\$207.80, BUY)
- PDD Holdings Inc. (PDD: \$88.56, BUY)
- Tencent Holdings Ltd. (700 HK: HK\$475.20, BUY)
- Trip.com Group (TCOM: \$47.01, BUY)
- Vipshop Holdings Limited (VIPS: \$15.34, BUY)
- Weibo Corporation (WB: \$8.15, BUY)
- Xiaomi Corporation (1810 HK: HK\$28.78, UNDERPERFORM)
- Z.AI (2513 HK: HK\$987.50, HOLD)

Distribution of Ratings

			IB Serv./Past12 Mos.		JIL Mkt Serv./Past12 Mos.	
	Count	Percent	Count	Percent	Count	Percent
BUY	2229	62.45%	392	17.59%	114	5.11%
HOLD	1180	33.06%	92	7.80%	14	1.19%
UNDERPERFORM	160	4.48%	0	0.00%	0	0.00%

Other Important Disclosure

Other Important Disclosures

Jefferies does business and seeks to do business with companies covered in its research reports, and expects to receive or intends to seek compensation for investment banking services among other activities from such companies. As a result, investors should be aware that Jefferies may have a conflict of interest that could affect the objectivity of this report. Investors should consider this report as only a single factor in making their investment decision.

Jefferies Equity Research refers to research reports produced by analysts employed by one of the following Jefferies Financial Group Inc. ("Jefferies") companies:

United States: Jefferies LLC which is an SEC registered broker-dealer and a member of FINRA (and distributed by Jefferies Research Services, LLC, an SEC registered Investment Adviser, to clients paying separately for such research).

Canada: Jefferies Securities Inc., which is an investment dealer registered in each of the thirteen Canadian jurisdictions and a dealer member of the Canadian Investment Regulatory Organization, including research reports produced jointly by Jefferies Securities Inc. and another Jefferies entity (and distributed by Jefferies Securities Inc.).

Where Jefferies Securities Inc. distributes research reports produced by Jefferies LLC, Jefferies International Limited, Jefferies Japan Company Limited, or Jefferies India Private Limited, you are advised that each of Jefferies LLC, Jefferies International Limited, Jefferies Japan Company Limited, and Jefferies India Private Limited operates as a dealer in your jurisdiction under an exemption from the dealer registration requirements contained in National Instrument 31-103 Registration Requirements, Exemptions and Ongoing Registrant Obligations and, as such, each of Jefferies LLC, Jefferies International Limited, Jefferies Japan Company Limited, and Jefferies India Private Limited is not required to be and is not a registered dealer or adviser in your jurisdiction. You are advised that where Jefferies LLC or Jefferies International Limited prepared this research report, it was not prepared in accordance with Canadian disclosure requirements relating to research reports in Canada.

United Kingdom: Jefferies International Limited, which is authorized and regulated by the Financial Conduct Authority; registered in England and Wales No. 1978621; registered office: 100 Bishopsgate, London EC2N 4JL; telephone +44 (0)20 7029 8000; facsimile +44 (0)20 7029 8010.

Germany: Jefferies GmbH, which is authorized and regulated by the Bundesanstalt fuer Finanzdienstleistungsaufsicht, BaFin-ID: 10150151; registered office: Bockenheimer Landstr. 24, 60323 Frankfurt a.M., Germany; telephone: +49 (0) 69 719 1870

Hong Kong: Jefferies Hong Kong Limited, which is licensed by the Securities and Futures Commission of Hong Kong with CE number ATS546; located at Level 26, Two International Finance Center, 8 Finance Street, Central, Hong Kong; telephone: +852 3743 8000.

Singapore: Jefferies Singapore Limited, which is licensed by the Monetary Authority of Singapore; located at 10 Collyer Quay #41-01, Ocean Financial Centre, Singapore 049315, telephone: +65 6551 3950.

Japan: Jefferies Japan Company Limited, which is a securities company registered by the Financial Services Agency of Japan and is a member of the Japan Securities Dealers Association; located at Tokyo Midtown Hibiya 30F Hibiya Mitsui Tower, 1-1-2 Yuraku-cho, Chiyoda-ku, Tokyo 100-0006; telephone +813 5251 6100; facsimile +813 5251 6101.

India: Jefferies India Private Limited (CIN - U74140MH2007PTC200509), licensed by the Securities and Exchange Board of India for: Stock Broker (NSE & BSE) INZ000243033, Research Analyst INH000000701 and Merchant Banker INM000011443, located at Level 16, Express Towers, Nariman Point, Mumbai 400 021, India; Tel +91 22 4356 6000. Compliance Officer name: Sanjay Pai, Tel No: +91 22 42246150, Email: spai@jefferies.com, Grievance officer name: Sanjay Pai, Tel no. +91 22 42246150, Email: compliance_india@jefferies.com. Registration granted by SEBI and certification from NISM in no way guarantee performance of the intermediary or provide any assurance of returns to investors.

Australia: Jefferies (Australia) Pty Limited (ACN 623 059 898), which holds an Australian financial services license (AFSL 504712) and is located at Level 20, 60 Martin Place, Sydney NSW 2000; telephone +61 2 9364 2800.

Dubai: Jefferies International Limited, Dubai branch, which is licensed by the Dubai Financial Services Authority (DFSA Reference Number F007325); registered office Unit L31-06, L31-07, Level 31, ICD Brookfield Pace, DIFC, PO Box 121208, Dubai, UAE.

This report was prepared by personnel who are associated with Jefferies (Jefferies Securities Inc., Jefferies International Limited, Jefferies GmbH, Jefferies Hong Kong Limited, Jefferies Singapore Limited, Jefferies Japan Company Limited, Jefferies India Private Limited), and Jefferies (Australia) Pty Ltd; or by personnel who are associated with both Jefferies LLC and Jefferies Research Services LLC ("JRS"). Jefferies LLC is a US registered broker-dealer and is affiliated with JRS, which is a US registered investment adviser. JRS does not create tailored or personalized research and all research provided by JRS is impersonal. If you are paying separately for this research, it is being provided to you by JRS. Otherwise, it is being provided by Jefferies LLC. Jefferies LLC, JRS, and their affiliates are collectively referred to below as "Jefferies". Jefferies may seek to do business with companies covered in this research report. As a result, investors should be aware that Jefferies may have a conflict of interest that could affect the objectivity of this report. Investors should consider this report as only one of many factors in making their investment decisions. Specific conflict of interest and other disclosures that are required by FINRA, the Canadian Investment Regulatory Organization and other rules are set forth in this disclosure section.

If you are receiving this report from a non-US Jefferies entity, please note the following: Unless prohibited by the provisions of Regulation S of the U.S. Securities Act of 1933, as amended, this material is distributed in the United States by Jefferies LLC, which accepts responsibility for its contents in accordance with the provisions of Rule 15a-6 under the US Securities Exchange Act of 1934, as amended. Transactions by or on behalf of any US person may only be effected through Jefferies LLC. In the United Kingdom and European Economic Area this report is issued and/or approved for distribution by Jefferies International Limited ("JIL") and/or Jefferies GmbH and is intended for use only by persons who have, or have been assessed as having, suitable professional experience and expertise, or by persons to whom it can be otherwise lawfully distributed. Jefferies LLC, JIL, Jefferies GmbH and their affiliates, may make a market or provide liquidity in the financial instruments referred to in this report; and where they do make a market, such activity is disclosed specifically in this report under "company specific disclosures".

For Canadian investors, this material is intended for use only by professional or institutional investors. None of the investments or investment services mentioned or described herein is available to other persons or to anyone in Canada who is not a "permitted client" as defined by National Instrument 31-103 Registration Requirements, Exemptions and Ongoing Registrant Obligations, as applicable. This research report is a general discussion of the merits and risks of a security or securities only, and is not in any way meant to be tailored to the needs and circumstances of any recipient. The information contained

herein is not, and under no circumstances is to be construed as, an offer to sell securities described herein, or solicitation of an offer to buy securities described herein, in Canada or any province or territory thereof. Any offer or sale of the securities described herein in Canada will be made only under an exemption from the requirements to file a prospectus with the relevant Canadian securities regulators, if applicable, and only by a dealer properly registered under applicable securities laws or, alternatively, pursuant to an exemption from the dealer registration requirement in the relevant province or territory of Canada in which such offer or sale is made. The information contained herein is under no circumstances to be construed as investment advice in any province or territory of Canada and is not tailored to the needs of the recipient. To the extent that the information contained herein references securities of an issuer incorporated, formed or created under the laws of Canada or a province or territory of Canada, any trades in such securities must be conducted through a dealer registered in Canada. No securities commission or similar regulatory authority in Canada has reviewed or in any way passed judgment upon this research report, the information contained herein or the merits of the securities described herein, and any representation to the contrary is an offence. In Singapore, Jefferies Singapore Limited ("JSL") is regulated by the Monetary Authority of Singapore. For investors in the Republic of Singapore, where this material is prepared and issued by a Jefferies affiliate outside of Singapore, it is distributed by JSL pursuant to Regulation 32C of the Financial Advisers Regulations. The material contained in this document is intended solely for accredited, expert or institutional investors, as defined under the Securities and Futures Act 2001 (Singapore). If there are any matters arising from, or in connection with this material, please contact JSL, located at 80 Raffles Place #15-20, UOB Plaza 2, Singapore 048624, telephone: +65 6551 3950. In Dubai, this material is issued and distributed by Jefferies International Limited, Dubai branch, and is intended solely for Professional Clients and should not be distributed to, or relied upon by, Retail Clients (as defined by DFSA). A distribution of ratings in percentage terms in each sector covered is available upon request from your sales representative. In Japan, this material is issued and distributed by Jefferies Japan Company Limited to institutional investors only. In Hong Kong, this report is issued and approved by Jefferies Hong Kong Limited and is intended for use only by professional investors as defined in the Hong Kong Securities and Futures Ordinance and its subsidiary legislation. In the Republic of China (Taiwan), this report should not be distributed. The research in relation to this report is conducted outside the People's Republic of China ("PRC"). This report does not constitute an offer to sell or the solicitation of an offer to buy any securities in the PRC. PRC investors shall have the relevant qualifications to invest in such securities and shall be responsible for obtaining all relevant approvals, licenses, verifications and/or registrations from the relevant governmental authorities themselves. In India, this report is made available by Jefferies India Private Limited. In Australia, this report is issued and/or approved for distribution by, or on behalf of, Jefferies (Australia) Securities Pty Ltd (ACN 610 977 074), which holds an Australian financial services license (AFSL 487263). It is directed solely at wholesale clients within the meaning of the Corporations Act 2001 (Cth) of Australia (the "Corporations Act"), in connection with their consideration of any investment or investment service that is the subject of this report. This report may contain general financial product advice. Where this report refers to a particular financial product, you should obtain a copy of the relevant product disclosure statement or offer document before making any decision in relation to the product. Recipients of this document in any other jurisdictions should inform themselves about and observe any applicable legal requirements in relation to the receipt of this document.

This report is not an offer or solicitation of an offer to buy or sell any security or derivative instrument, or to make any investment. Any opinion or estimate constitutes the preparer's best judgment as of the date of preparation, and is subject to change without notice. Jefferies assumes no obligation to maintain or update this report based on subsequent information and events. Jefferies, and their respective officers, directors, and employees, may have long or short positions in, or may buy or sell any of the securities, derivative instruments or other investments mentioned or described herein, either as agent or as principal for their own account. This material is provided solely for informational purposes and is not tailored to any recipient, and is not based on, and does not take into account, the particular investment objectives, portfolio holdings, strategy, financial situation, or needs of any recipient. As such, any advice or recommendation in this report may not be suitable for a particular recipient. Jefferies assumes recipients of this report are capable of evaluating the information contained herein and of exercising independent judgment. A recipient of this report should not make any investment decision without first considering whether any advice or recommendation in this report is suitable for the recipient based on the recipient's particular circumstances and, if appropriate or otherwise needed, seeking professional advice, including tax advice. Jefferies does not perform any suitability or other analysis to check whether an investment decision made by the recipient based on this report is consistent with a recipient's investment objectives, portfolio holdings, strategy, financial situation, or needs.

By providing this report, neither JRS nor any other Jefferies entity accepts any authority, discretion, or control over the management of the recipient's assets. Any action taken by the recipient of this report, based on the information in the report, is at the recipient's sole judgment and risk. The recipient must perform his or her own independent review of any prospective investment. If the recipient uses the services of Jefferies LLC (or other affiliated broker-dealers), in connection with a purchase or sale of a security that is a subject of these materials, such broker-dealer may act as principal for its own accounts or as agent for another person. Only JRS is registered with the SEC as an investment adviser, and therefore neither Jefferies LLC nor any other Jefferies affiliate has any fiduciary duty in connection with distribution of these reports.

The price and value of the investments referred to herein and the income from them may fluctuate. Past performance is not a guide to future performance, future returns are not guaranteed, and a loss of original capital may occur. Fluctuations in exchange rates could have adverse effects on the value or price of, or income derived from, certain investments.

This report may contain forward looking statements that may be affected by inaccurate assumptions or by known or unknown risks, uncertainties, and other important factors. As a result, the actual results, events, performance or achievements of the financial product may be materially different from those expressed or implied in such statements.

This report has been prepared independently of any issuer of securities mentioned herein and not as agent of any issuer of securities. No Equity Research personnel have authority whatsoever to make any representations or warranty on behalf of the issuer(s). Any comments or statements made herein are those of the Jefferies entity producing this report and may differ from the views of other Jefferies entities.

This report may contain information obtained from third parties, including ratings from credit ratings agencies such as Standard & Poor's, and information derived from third-party or proprietary generative artificial intelligence (Gen AI) models. Jefferies does not guarantee the accuracy, completeness, timeliness or availability of this information, and is not responsible for any errors or omissions (negligent or otherwise), regardless of the cause, or for the results obtained from the use of such content. Neither Jefferies nor any third-party content providers, including providers of Gen AI models, give any express or implied warranties, including, but not limited to, any warranties of merchantability or fitness for a particular purpose or use. Neither Jefferies nor any third-party content provider shall be liable for any direct, indirect, incidental, exemplary, compensatory, punitive, special or consequential damages, costs, expenses, legal fees, or losses (including lost income or profits and opportunity costs) in connection with any use of their content, including ratings. Credit

ratings are statements of opinions and are not statements of fact or recommendations to purchase, hold or sell securities. They do not address the suitability of securities or the suitability of securities for investment purposes, and should not be relied on as investment advice. Reproduction and distribution of third party content in any form is prohibited except with the prior written permission of the related third party.

Jefferies research reports are disseminated and available electronically, and, in some cases, also in printed form. Electronic research is simultaneously made available to all clients. This report or any portion hereof may not be copied, reprinted, sold, or redistributed or disclosed by the recipient or any third party, by content scraping or extraction, automated processing, or any other form or means, without the prior written consent of Jefferies. Any unauthorized use is prohibited. Neither Jefferies nor any of its respective directors, officers or employees, is responsible for guaranteeing the financial success of any investment, or accepts any liability whatsoever for any direct, indirect or consequential damages or losses arising from any use of this report or its contents. Nothing herein shall be construed to waive any liability Jefferies has under applicable U.S. federal or state securities laws.

For Important Disclosure information relating to JRS, please see https://adviserinfo.sec.gov/IAPD/Content/Common/crd_iapd_Brochure.aspx?BRCHR_VRSN_ID=483878 and <https://adviserinfo.sec.gov/Firm/292142> or visit our website at <https://javatar.bluematrix.com/sellside/Disclosures.action>, or www.jefferies.com, or call 1.888.JEPPERIES.

© 2026 Jefferies